
1-Week Extension to Lab06, Lec06

Survey Review

About the course and critique, please.																										
3	2	2	2	2	1	1	2	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
3	2	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	3	2	1	2	3	2	1	2	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	3	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
8	3	2	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
7	1	1	1	1	1	1	2	1</td																		

Lectures

- Slides
 - More written notes
 - More organization
 - More example
 - Hands-on Practice
- Class
 - Chatbox
 - Q&A

Assignments

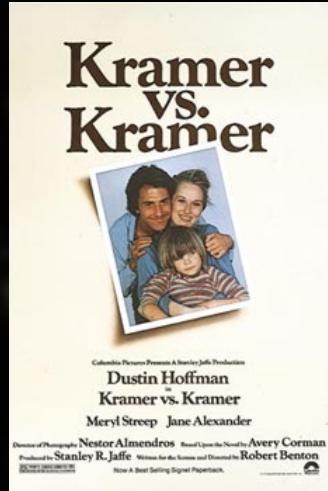
- Lecs
 - Better wording
 - More weights (time-benefit trade-off)
- Labs
 - More details
- Not Sync

Exam

- Not more than class duration
- Some questions not covered in class



Are robots becoming more human or humans becoming more robotic?



Googler vs. Googler

Googler: A Person Who Builds Google
Googler: A Person Who Works for Google
Googler: A Person Who Uses Google

Operating System Developer vs. Operating System Administrator Shell Developer vs. Bash Script Writer



Brian Jhan Fox



Linus Torvalds



USENIX: UNIX Users, New York, 1974

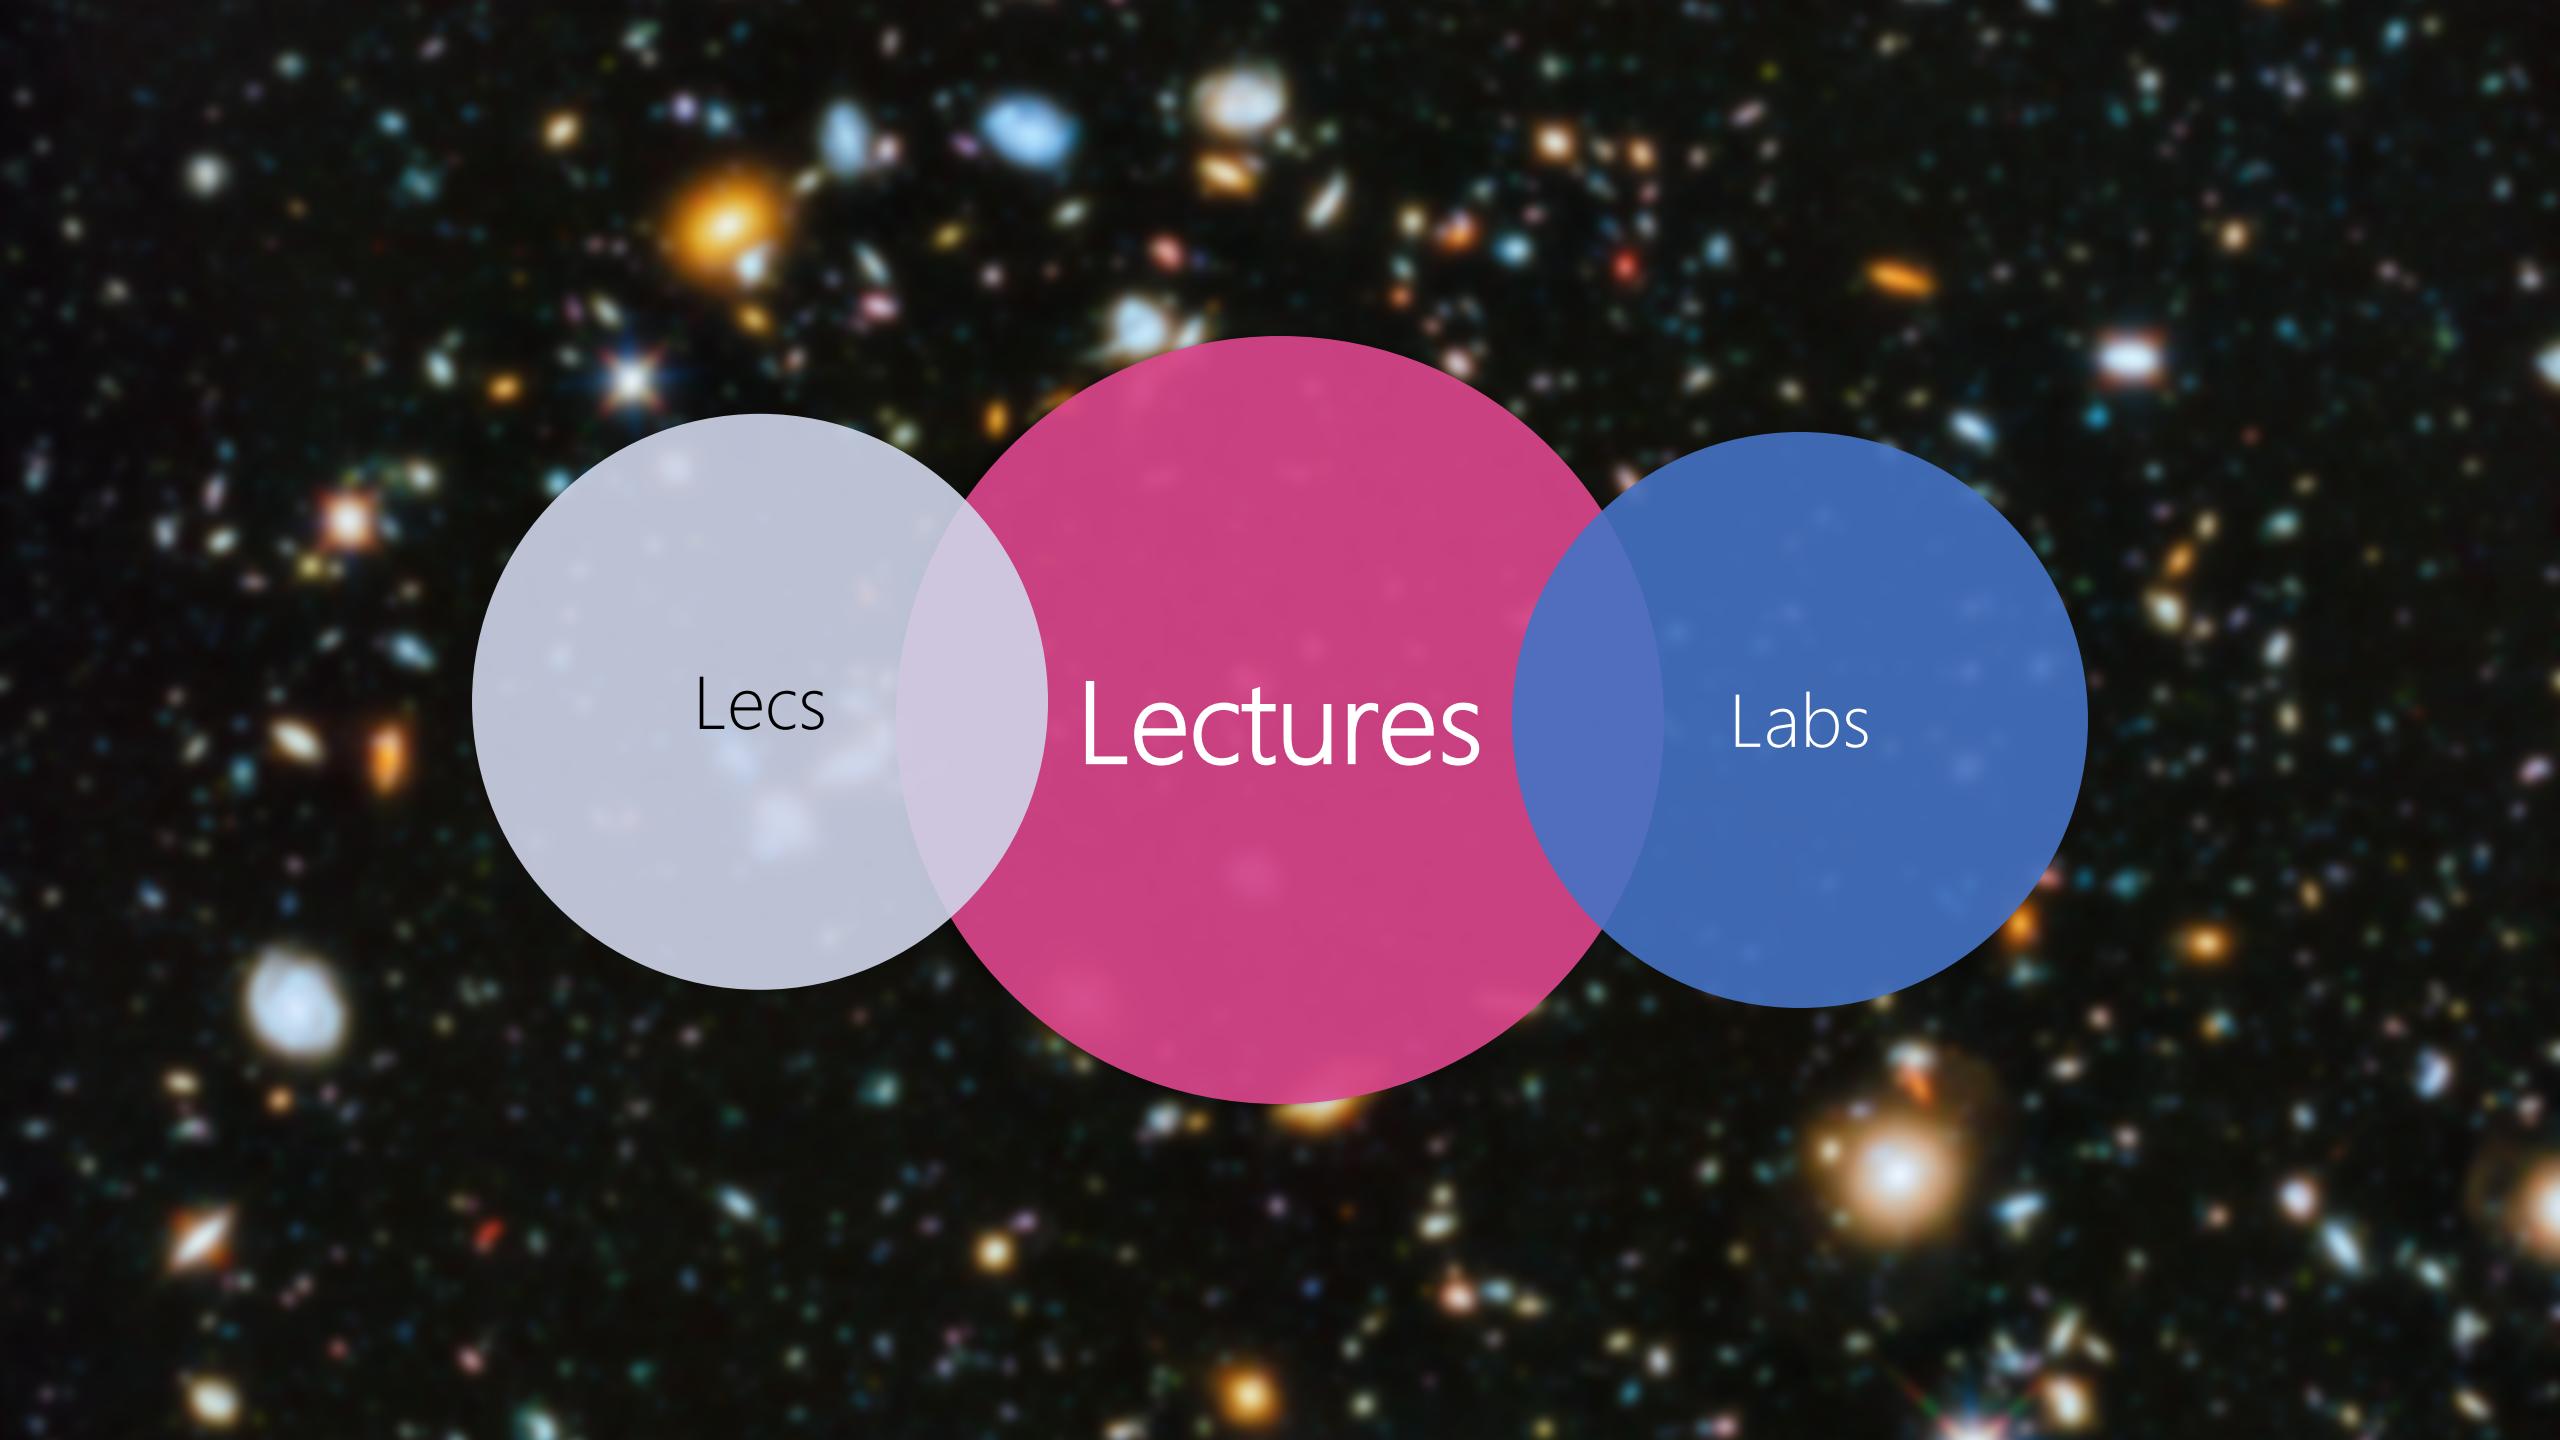
Lecture
Presentation
Dialog

vs.

vs.

vs.

Text
Book
Monolog



Lecs

Lectures

Labs

Lectures

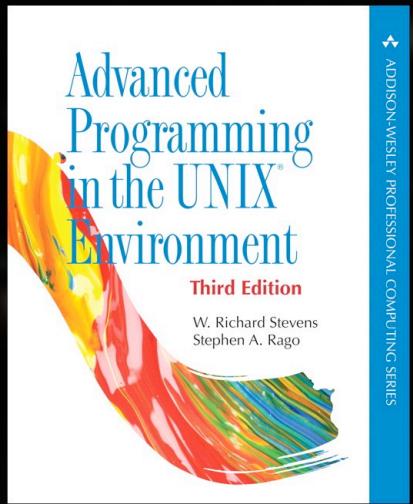
- Slides
 - More written notes → Yes but won't replace lectures!
 - More organization → Can't agree more.
 - More example → Sure. As we're getting toward coding
 - Hands-on Practice → No, please attend lab sections!
- Class
 - Chatbox → True. I used to have moderator.
 - Q&A → True. I try to put Q&A slides when moving from one topic to new one

Assignments

- Lecs
 - Better wording → True! I'll try to fix it.
 - More weights (time-benefit trade-off) → No.
 - Those scared of coding, start coding!
 - Others, this course is a technical course: Labs > Lecs
- Labs
 - More details → No. Attend the lab sections instead.
 - Not Sync → True. I'll try to fix it.

Exam

- Not more than class duration → True! I'll try to fix it.
- Some questions not covered in class → ?



Chapter 03: File I/O
Chapter 04: File I/O

Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

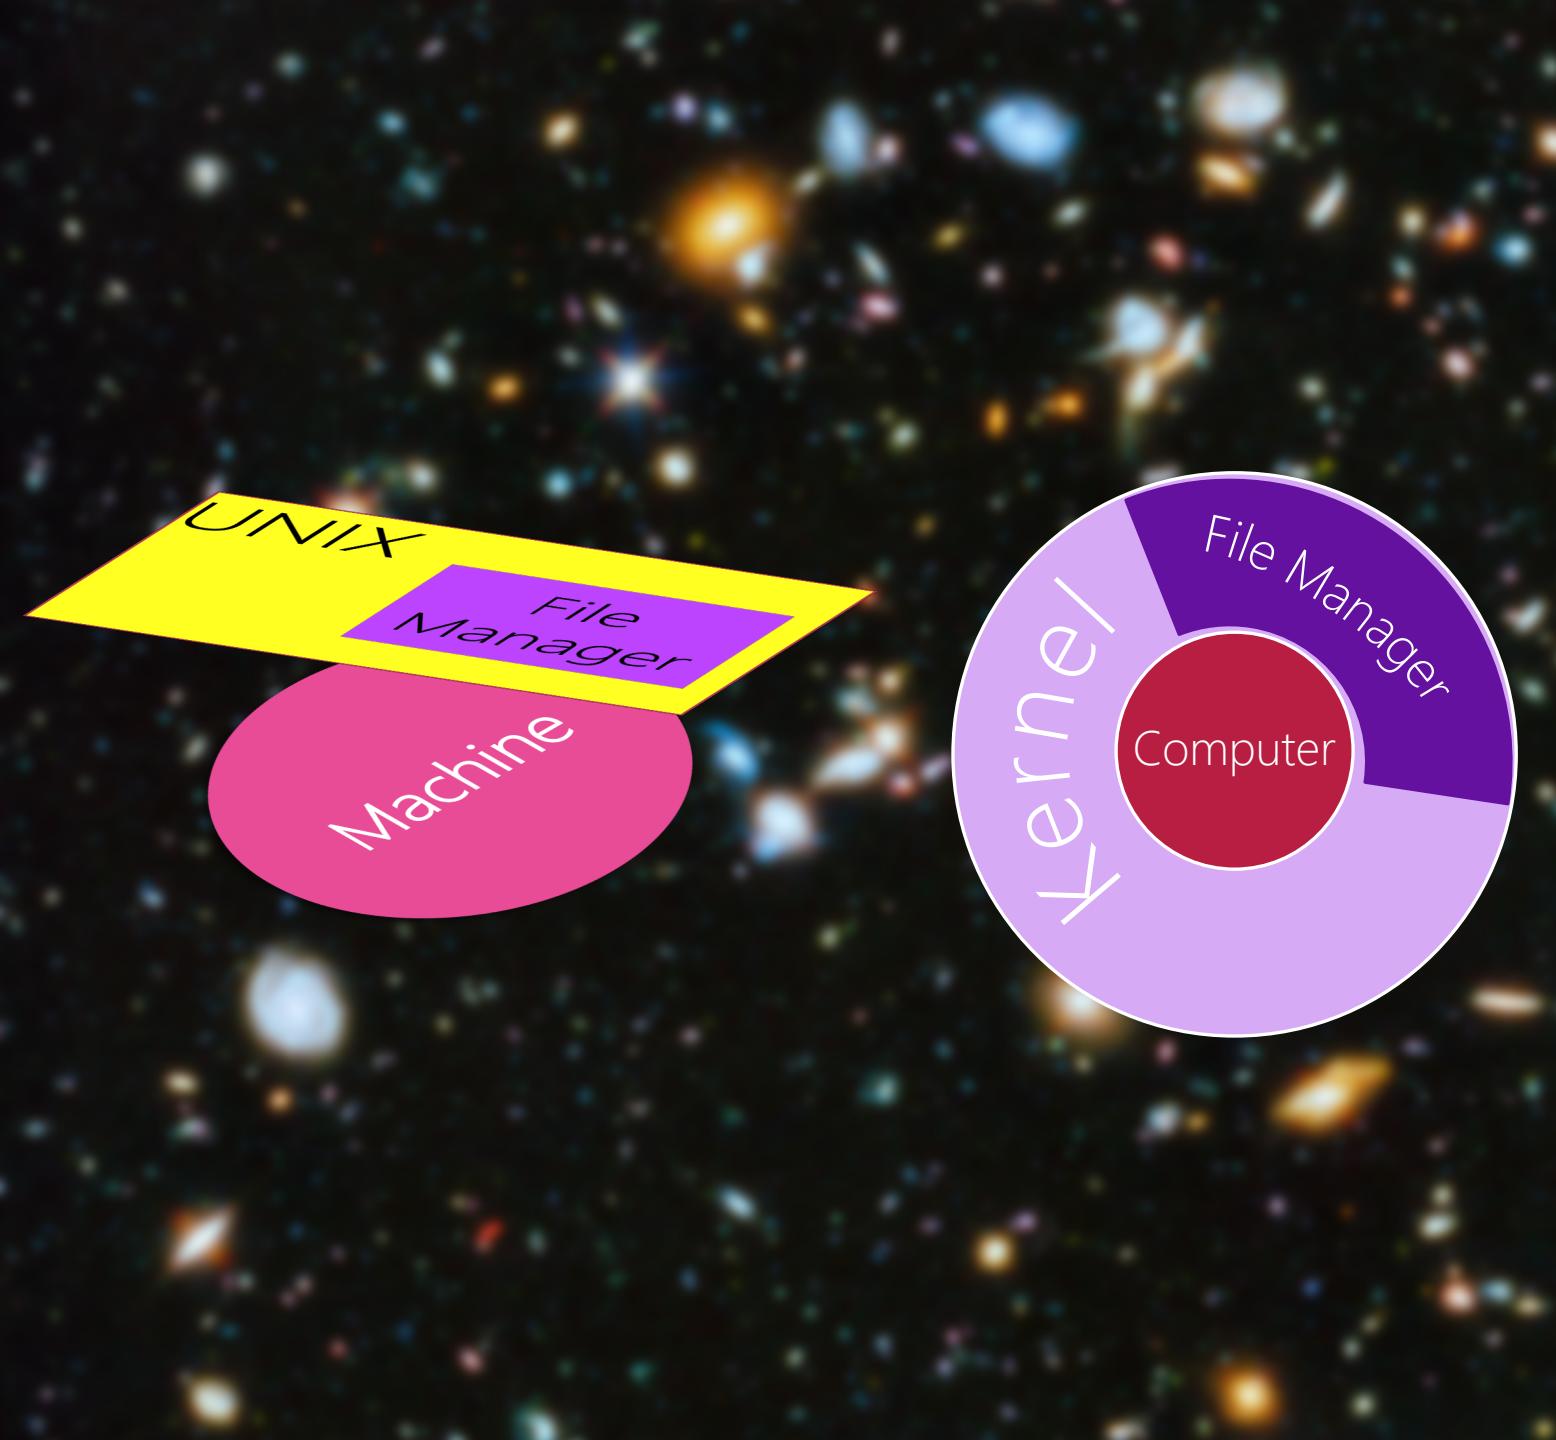
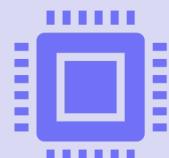
Kernel: File Manager

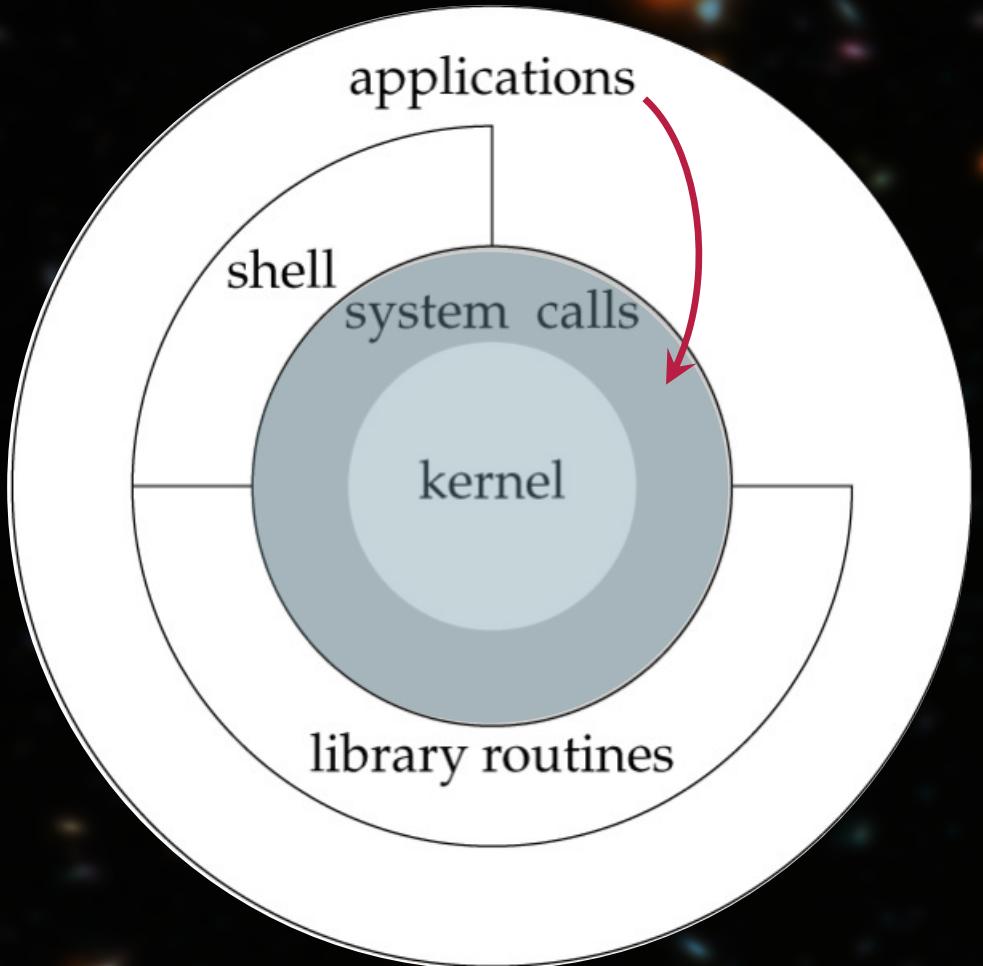
Kernel: Network Manager

Kernel: Process Manager



Processor





Header	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
<aio.h>	•	•	•	•	asynchronous I/O
<cpio.h>	•	•	•	•	cpio archive values
<dirent.h>	•	•	•	•	directory entries (Section 4.22)
<dlfcn.h>	•	•	•	•	dynamic linking
<fcntl.h>	•	•	•	•	file control (Section 3.14)
<fnmatch.h>	•	•	•	•	filename-matching types
<glob.h>	•	•	•	•	pathname pattern-matching and generation
<grp.h>	•	•	•	•	group file (Section 6.4)
<iconv.h>	•	•	•	•	codeset conversion utility
<langinfo.h>	•	•	•	•	language information constants
<monetary.h>	•	•	•	•	monetary types and functions
<netdb.h>	•	•	•	•	network database operations
<nl_types.h>	•	•	•	•	message catalogs
<poll.h>	•	•	•	•	poll function (Section 14.4.2)
<pthread.h>	•	•	•	•	threads (Chapters 11 and 12)
<pwd.h>	•	•	•	•	password file (Section 6.2)
<regex.h>	•	•	•	•	regular expressions
<sched.h>	•	•	•	•	execution scheduling
<semaphore.h>	•	•	•	•	semaphores
<strings.h>	•	•	•	•	string operations
<tar.h>	•	•	•	•	tar archive values
<termios.h>	•	•	•	•	terminal I/O (Chapter 18)
<unistd.h>	•	•	•	•	symbolic constants
<wordexp.h>	•	•	•	•	word-expansion definitions
<arpa/inet.h>	•	•	•	•	Internet definitions (Chapter 16)
<net/if.h>	•	•	•	•	socket local interfaces (Chapter 16)
<netinet/in.h>	•	•	•	•	Internet address family (Section 16.3)
<netinet/tcp.h>	•	•	•	•	Transmission Control Protocol definitions
<sys/mman.h>	•	•	•	•	memory management declarations
<sys/select.h>	•	•	•	•	select function (Section 14.4.1)
<sys/socket.h>	•	•	•	•	sockets interface (Chapter 16)
<sys/stat.h>	•	•	•	•	file status (Chapter 4)
<sys/statvfs.h>	•	•	•	•	file system information
<sys/times.h>	•	•	•	•	process times (Section 8.17)
<sys/types.h>	•	•	•	•	primitive system data types (Section 2.8)
<sys/un.h>	•	•	•	•	UNIX domain socket definitions (Section 17.2)
<sys/utsname.h>	•	•	•	•	system name (Section 6.9)
<sys/wait.h>	•	•	•	•	process control (Section 8.6)

```
fd = creat()
```

```
write(fd, sth)
```

```
close(fd)
```

```
fd = open(creat if not exists)
```

```
read(fd, sth)
```

```
write(fd, sth)
```

```
close(fd)
```

O_RDONLY	Open for reading only (the returned fd can only read)
O_WRONLY	Open for write only (the returned fd can only write like <code>creat()</code>)
O_RDWR	Open for reading and writing (the returned fd can do both read and write)
O_EXEC	Open for execute only (the returned fd execute)
O_SEARCH	Open for search only (for directories)

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
non-negative number (fd) if OK
-1 on error
```



You Want me on that Wall, You Need me on that Wall
<https://www.youtube.com/watch?v=xd1fs5nWol>
A Few Good Men - Aaron Sorkin

lseek

POSIX

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

Every opened file has one sentinel: `current_offset`
Measures the number of bytes from the beginning of the file.

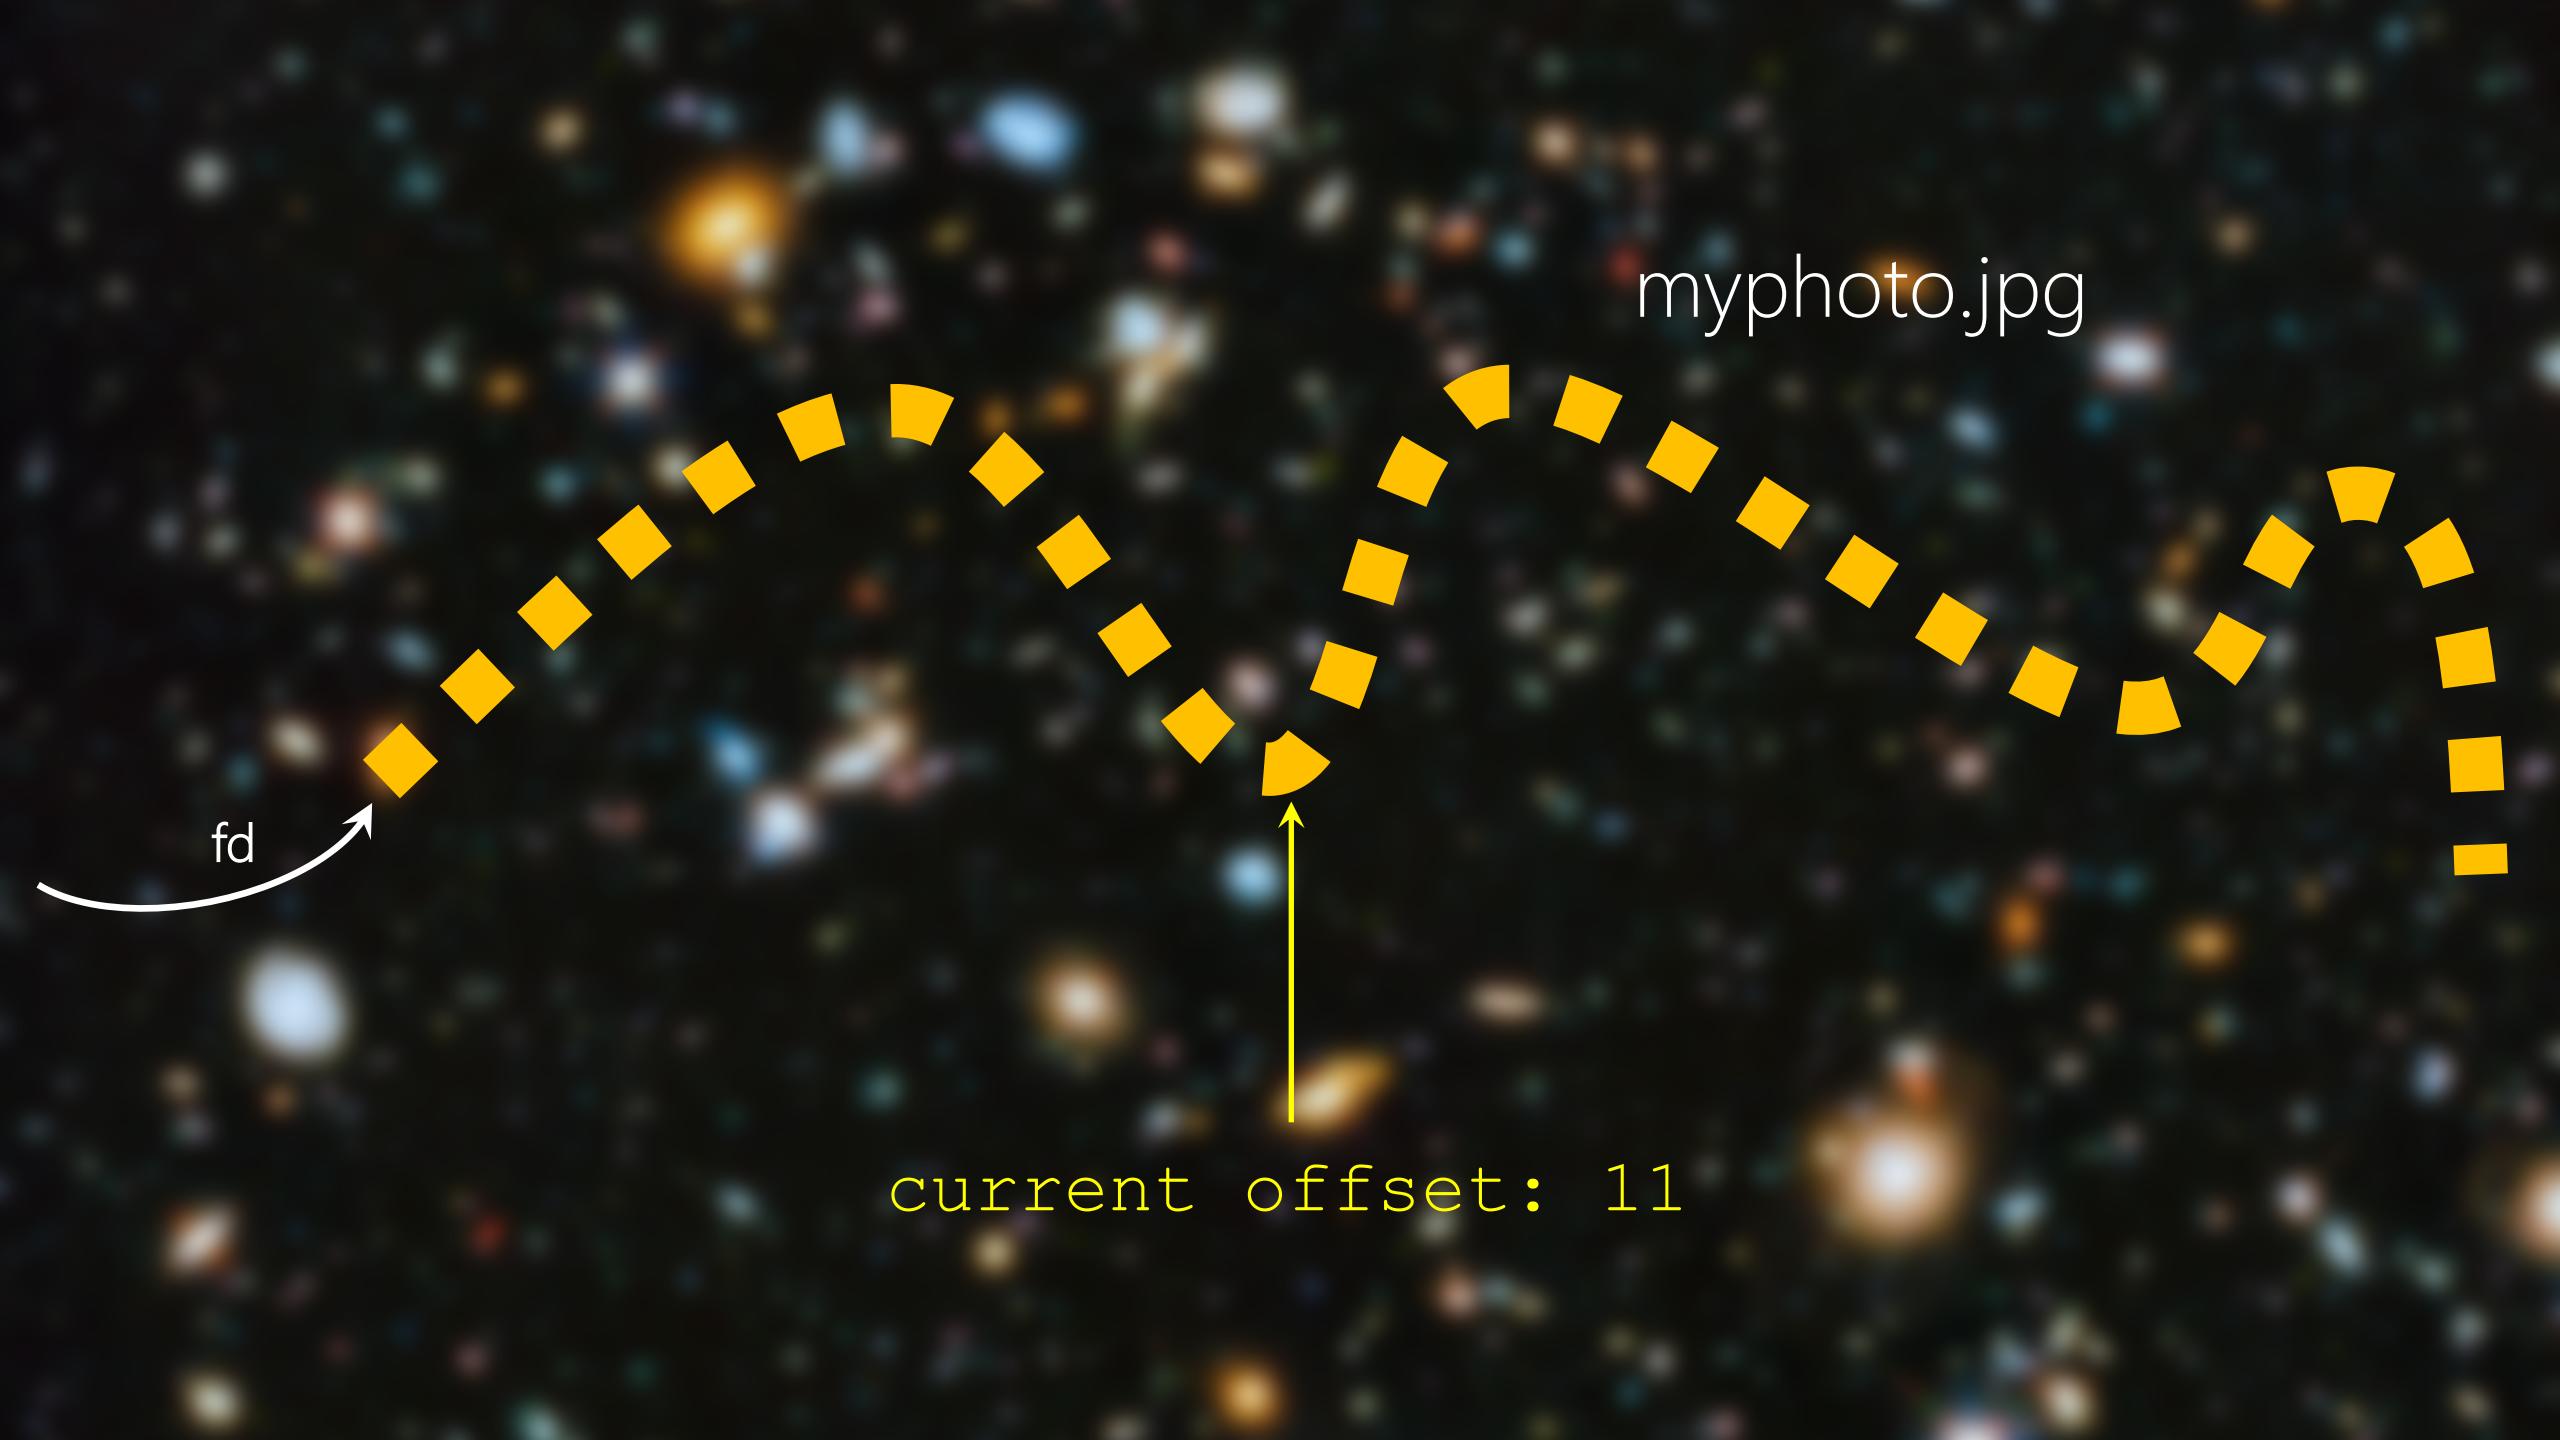
Every opened file has one sentinel: `current_offset`
Measures the number of bytes from the beginning of the file.

`creat()` or `open()` → 0

Every opened file has one sentinel: `current_offset`
Measures the number of bytes from the beginning of the file.

`read()` or `write()` → `++ actual` number of bytes read or written

`read()` or `write()` → always move forward



myphoto.jpg

fd



current offset: 11

lseek

How many bytes move the current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
```

file's new offset if OK (can be negative)
-1 on error

```
#include <sys/types.h>
typedef off_t signed long
```

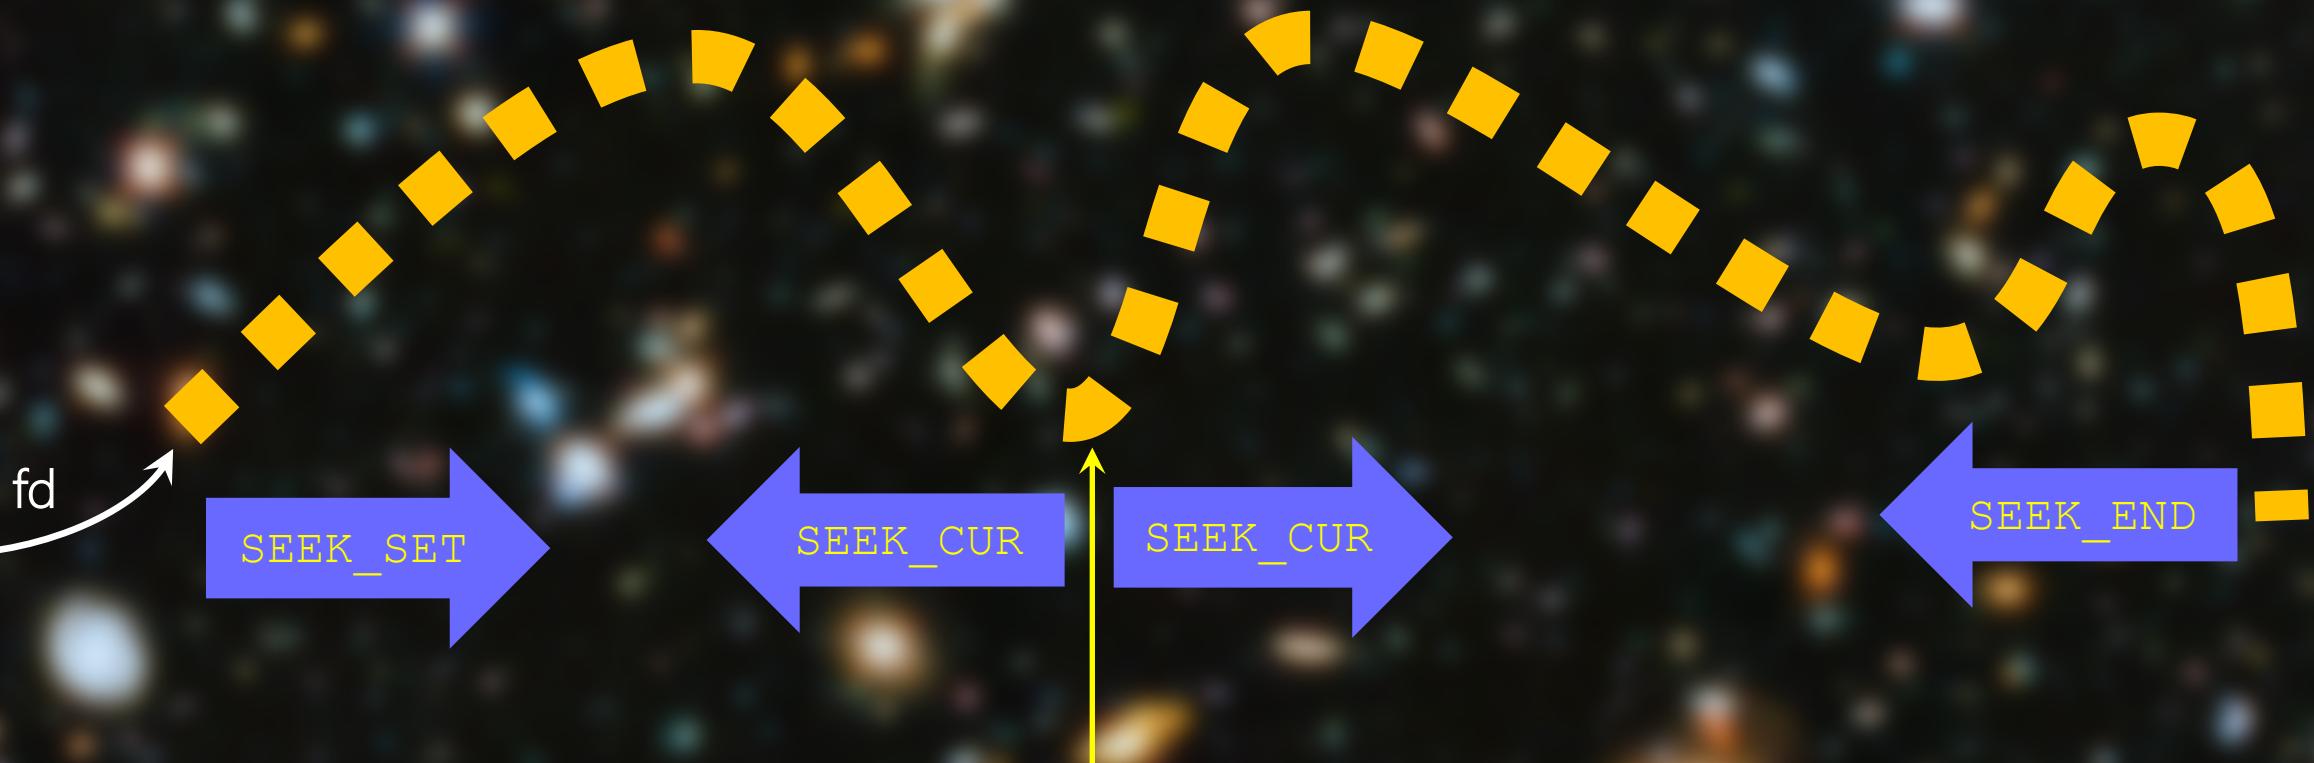
lseek

How many bytes move the current offset from what place or origin?

SEEK_SET, SEEK_CUR, SEEK_END

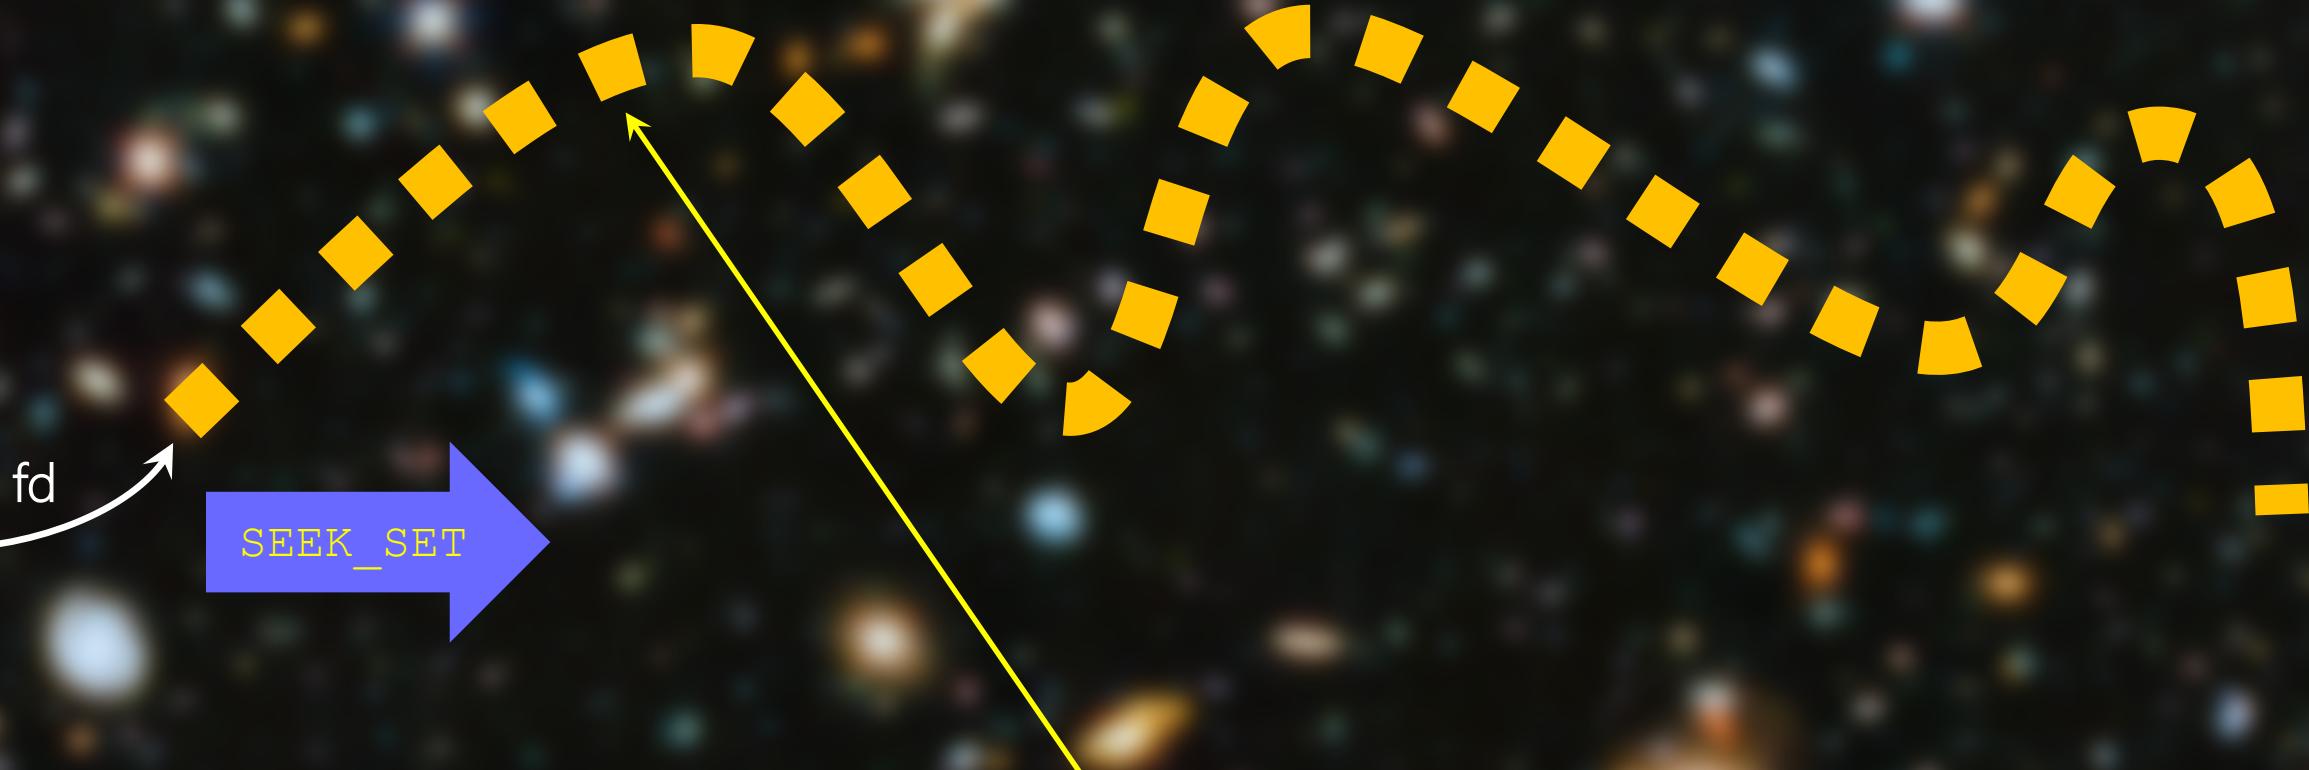
```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

myphoto.jpg



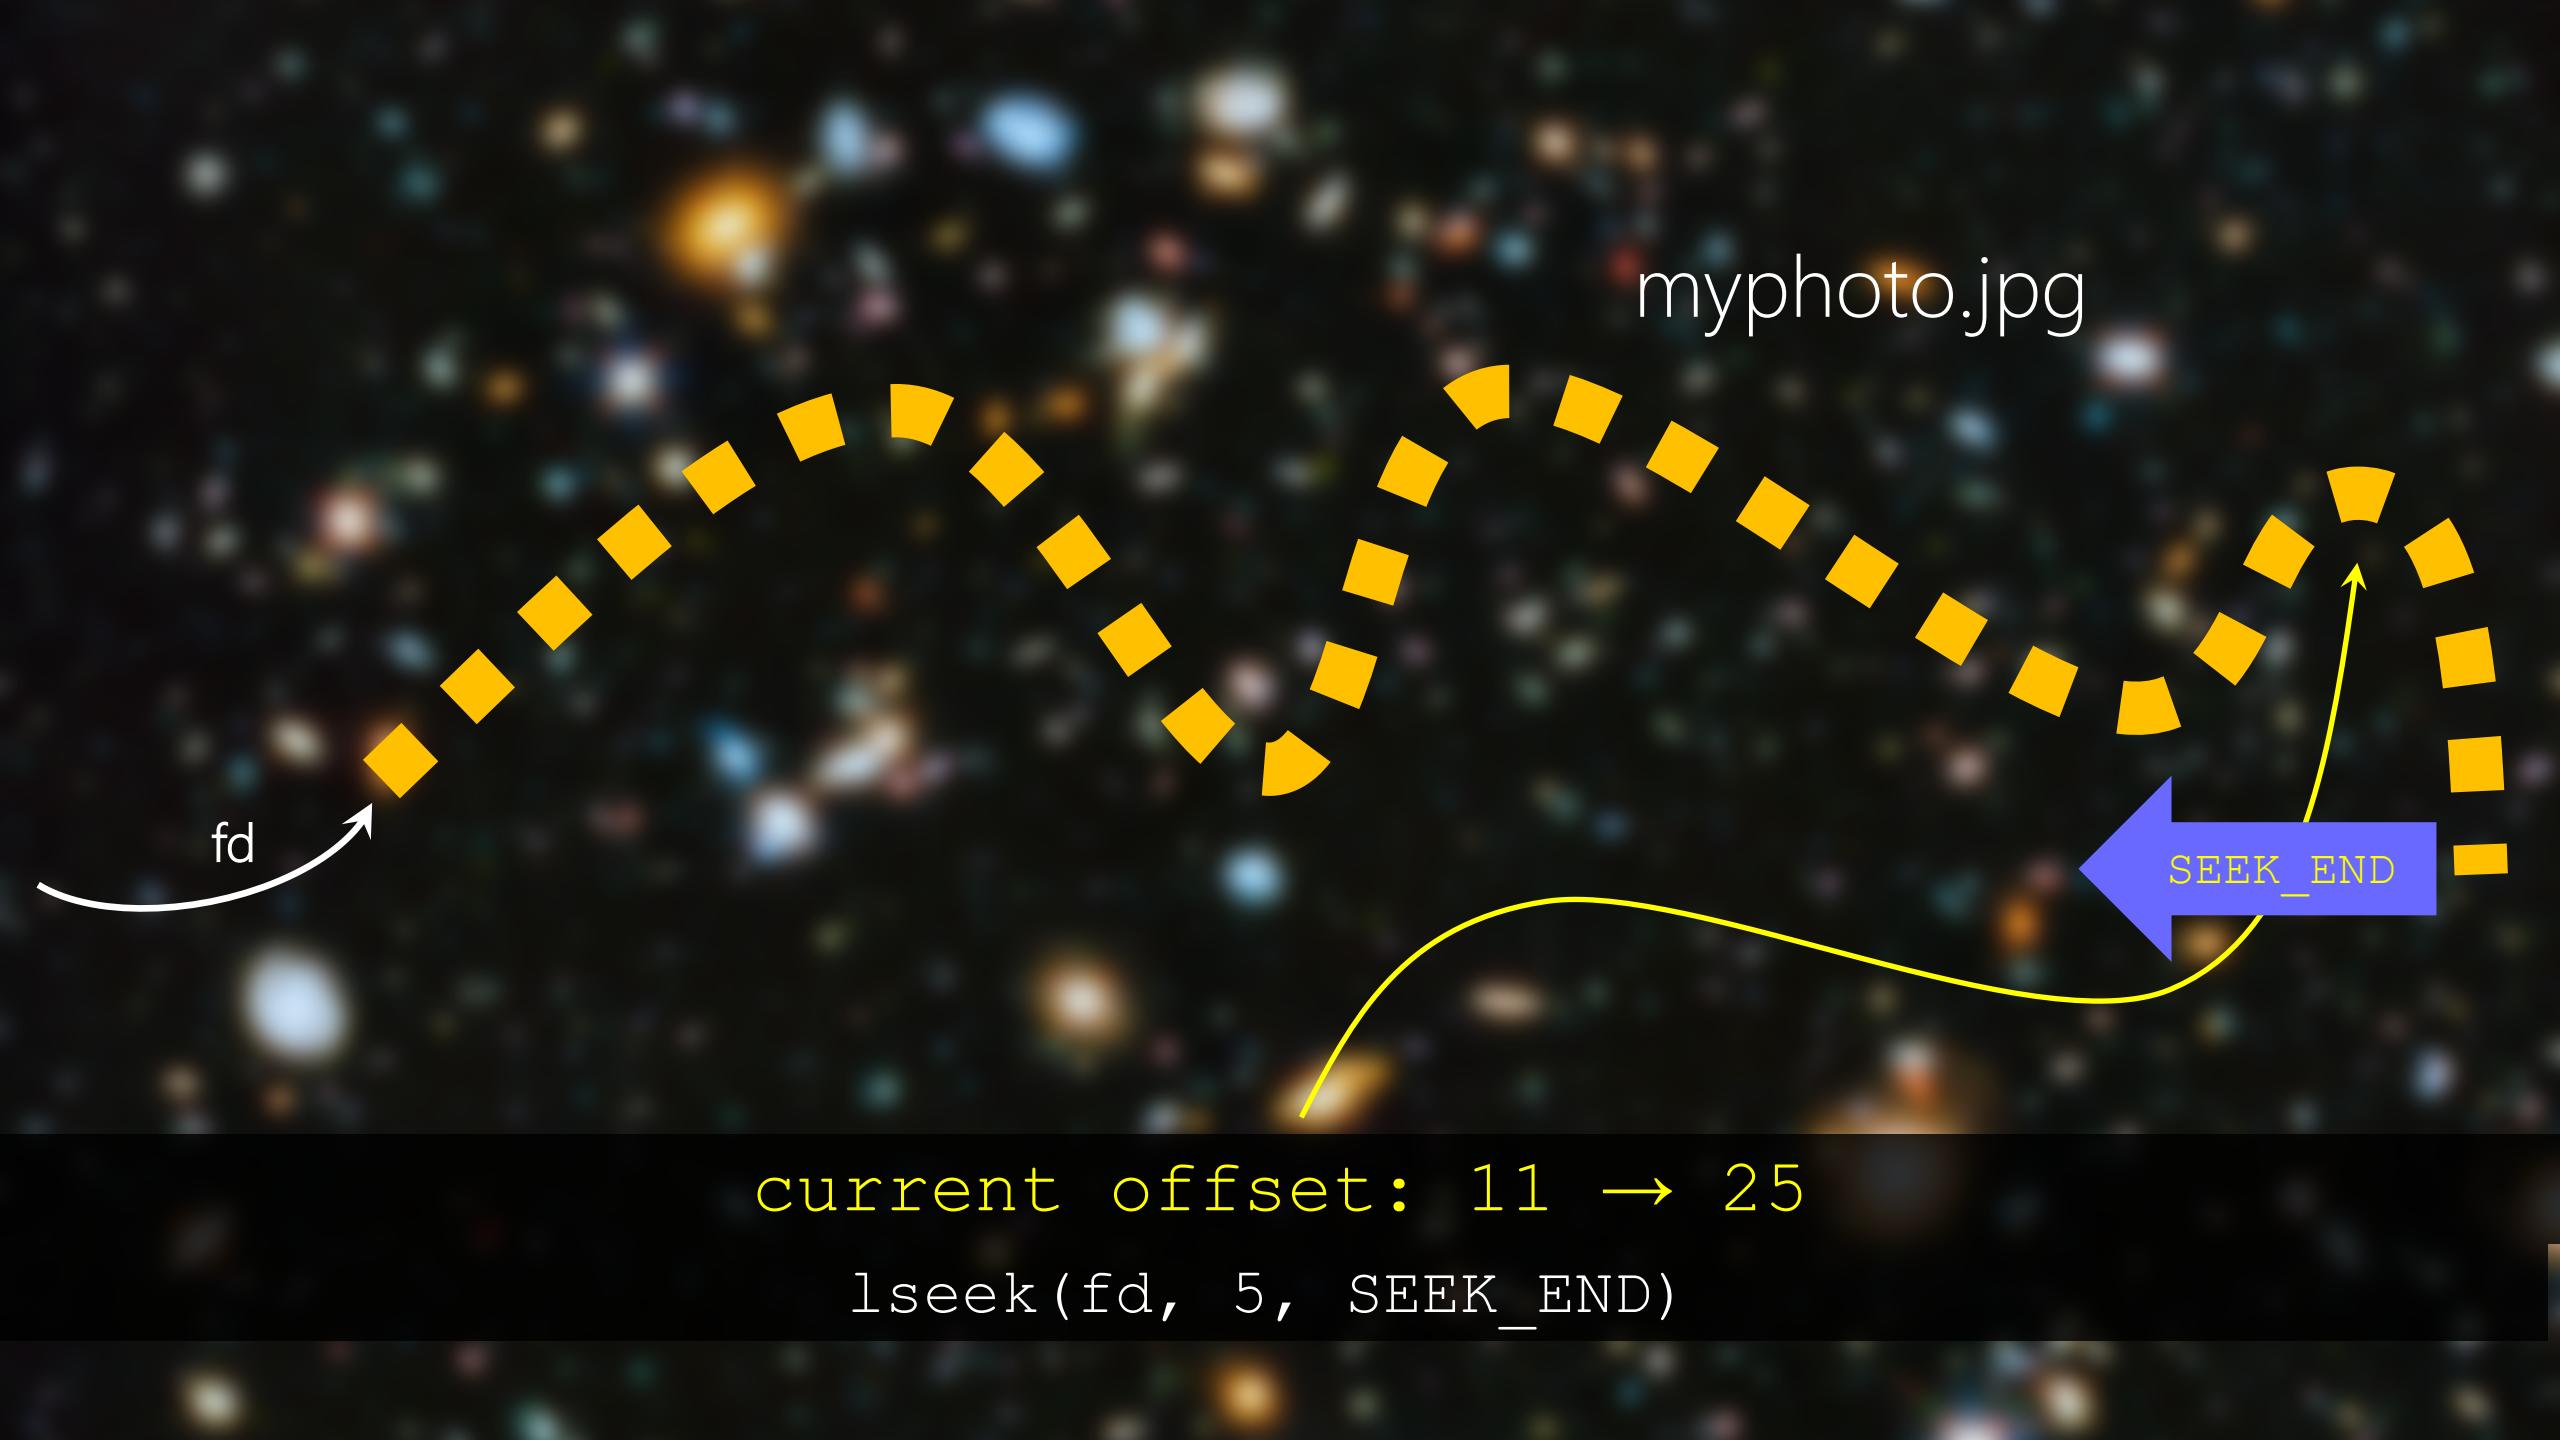
current offset: 11

myphoto.jpg



current offset: 11 → 5

lseek(fd, 5, SEEK_SET)



myphoto.jpg

The diagram illustrates a seek operation on a file named "myphoto.jpg". The file is represented by a dashed yellow arrow pointing from left to right, with several yellow diamond shapes along its path. A curved white arrow labeled "fd" points to the start of the arrow. A blue arrow labeled "SEEK_END" points from the right side of the arrow back towards the left, indicating a seek operation from the end of the file back towards the beginning. The background is a dark space-like texture with colorful stars.

fd

SEEK_END

current offset: 11 → 25

lseek(fd, 5, SEEK_END)

myphoto.jpg

The diagram illustrates a seek operation on a file named "myphoto.jpg". A dashed yellow arrow starts at the beginning of the file and points to a blue arrow labeled "SEEK_CUR". The blue arrow points to a specific byte position within the file. A curved white arrow labeled "fd" points to the start of the dashed arrow.

fd

SEEK_CUR

current offset: 11 → +5 = 16

lseek(fd, 5, SEEK_CUR)

myphoto.jpg



current offset: 11 → -5 = 6

`lseek(fd, -5, SEEK_CUR)`

lseek

New current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

lseek

How to know the value of current offset?

lseek

How to know the value of current offset?

```
long cur_offset;  
cur_offset = lseek(fd, 0, SEEK_CUR);
```

lseek

How to know the file can be seekable?

```
long cur_offset;
cur_offset = lseek(fd, 0, SEEK_CUR);
if (cur_offset == -1){
    /* either the fd is on a file which is not seekable
       or an error occurred */
}
```

```
hfani@charlie:~$ tty  
/dev/pts/12 ←  
hfani@charlie:~$ vi seekable.c
```

Is terminal device seekable?

```
hfani@charlie:~$ cc seekable.c -o seekable  
hfani@charlie:~$ ./seekable  
cannot seek! ←
```

Is terminal device seekable? No!

lseek

Is it possible for offset be greater than file size?



lseek

Is it possible for offset be greater than file size?
Yes! The next `write()` operation creates a hole of 0s.



```
hfani@charlie:~/Documents$ vi hole.c
#include <fcntl.h>
#include <unistd.h>
void main(void){
    int fd = open("./hole_test.txt", O_RDWR | O_CREAT, S_IRUSR | S_IWUSR); ← current offset = 0
    int cur_offset = lseek(fd, 10, SEEK_SET); ← move it 10 bytes ahead from start
    char buf[20] = "write after the hole.";
    write(fd, buf, 20);
}
```

```
hfani@charlie:~/Documents$ cc hole.c -o hole
hole.c: In function 'main':
hole.c:6:17: warning: initializer-string for array of 'char' is too long
  6 |     char buf[20] = "write after the hole.";

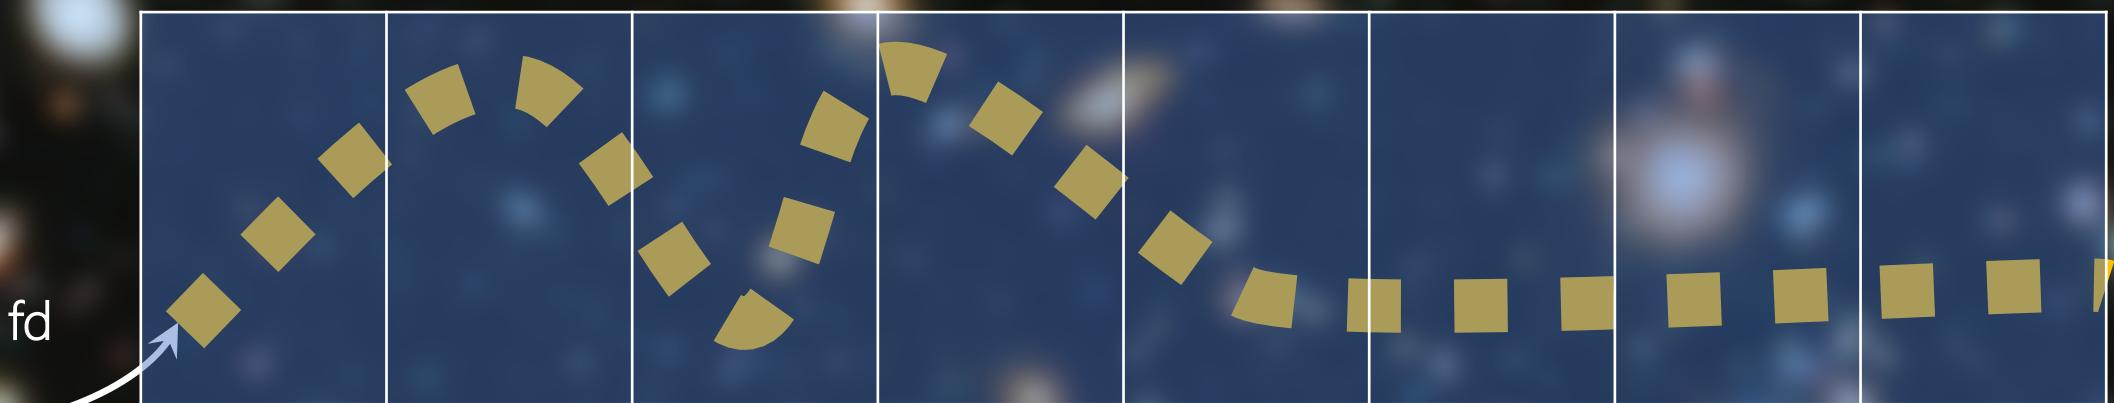
```

```
hfani@charlie:~$ hexdump hole_test.txt
00000000  0000 0000 0000 0000 7277 7469 2065
00000010  6661 6574 2072 6874 2065 6f68 656c
0000001e
```

```
hfani@charlie:~/Documents$ od -c hole test.txt
00000000  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  w  r  i  t  e
00000020  a  f  t  e  r          t  h  e          h  o  l  e
00000036
```

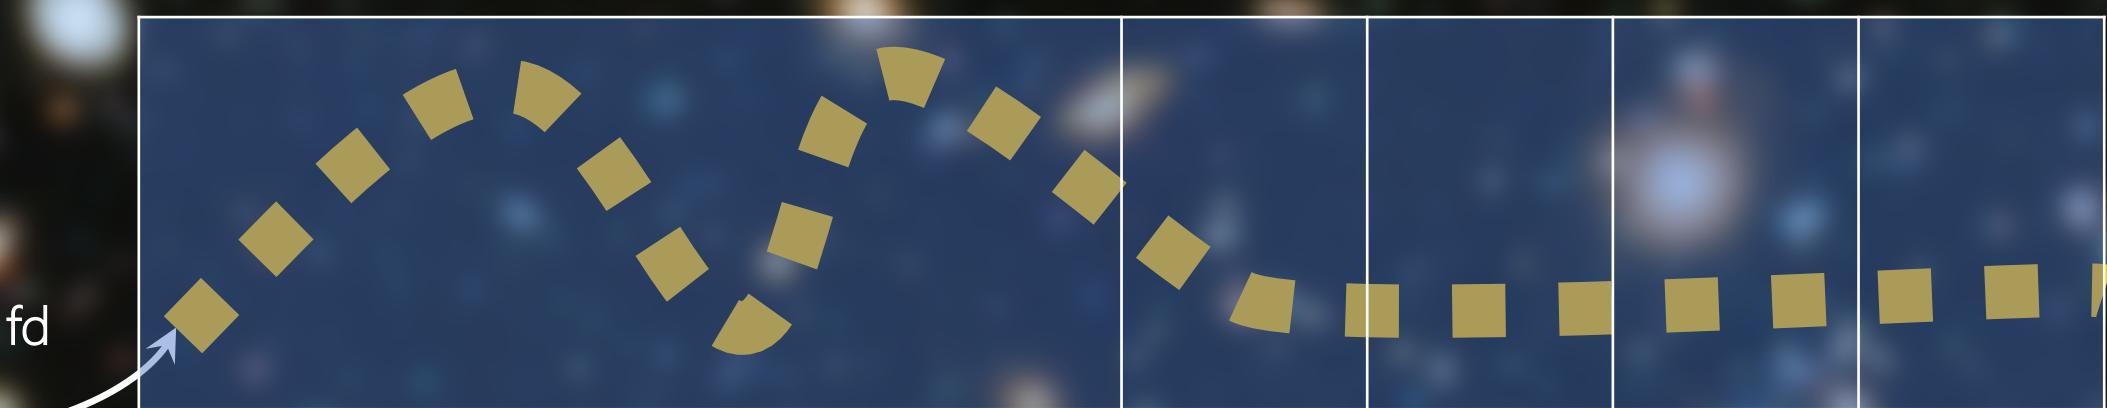
lseek

A Better User Case: Binary Search



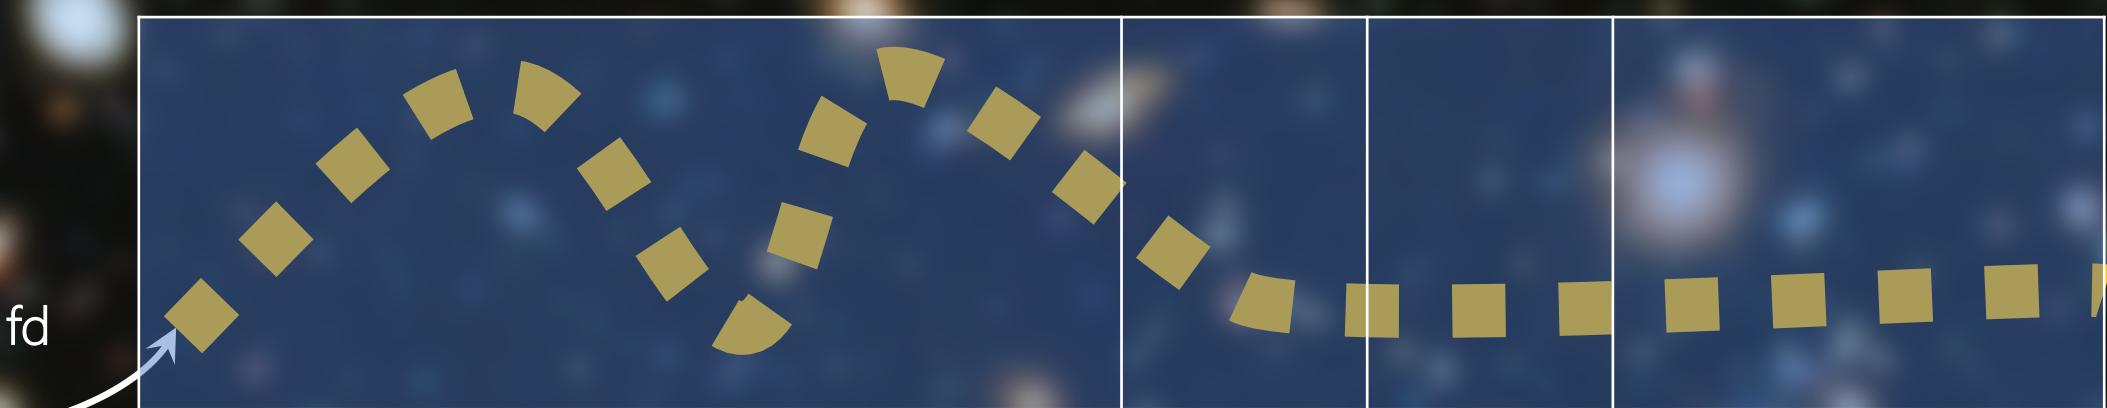
lseek

A Better User Case: Binary Search



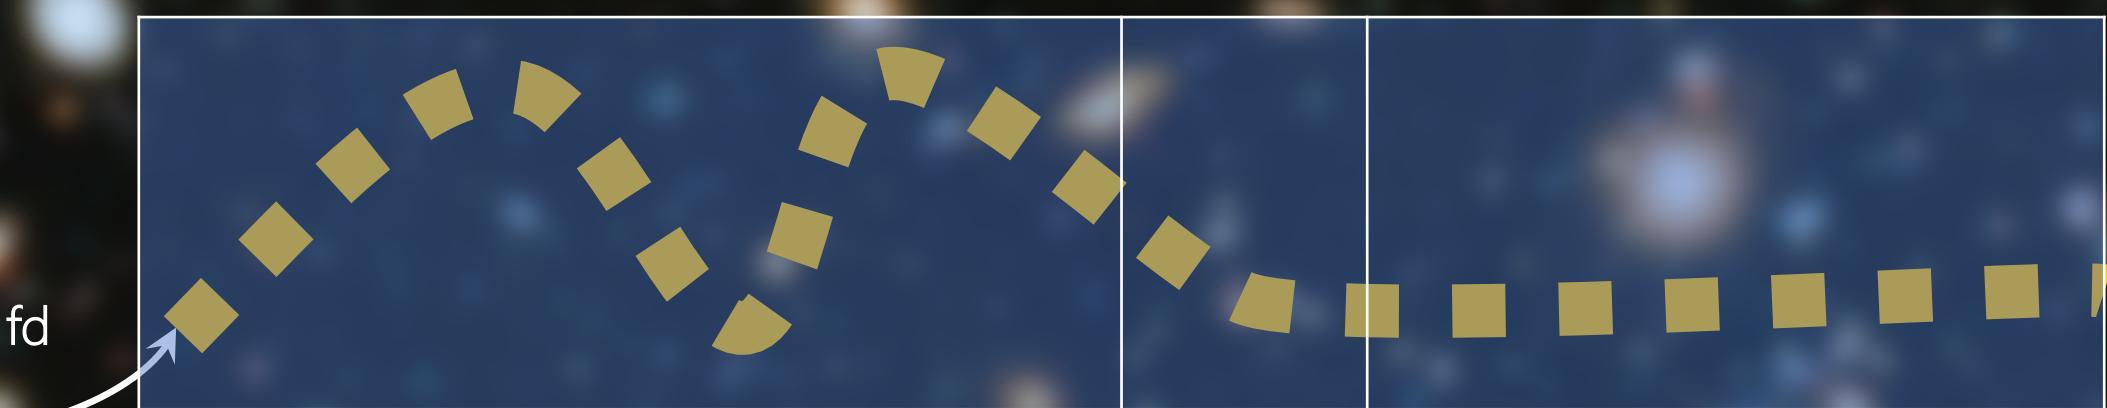
lseek

A Better User Case: Binary Search



lseek

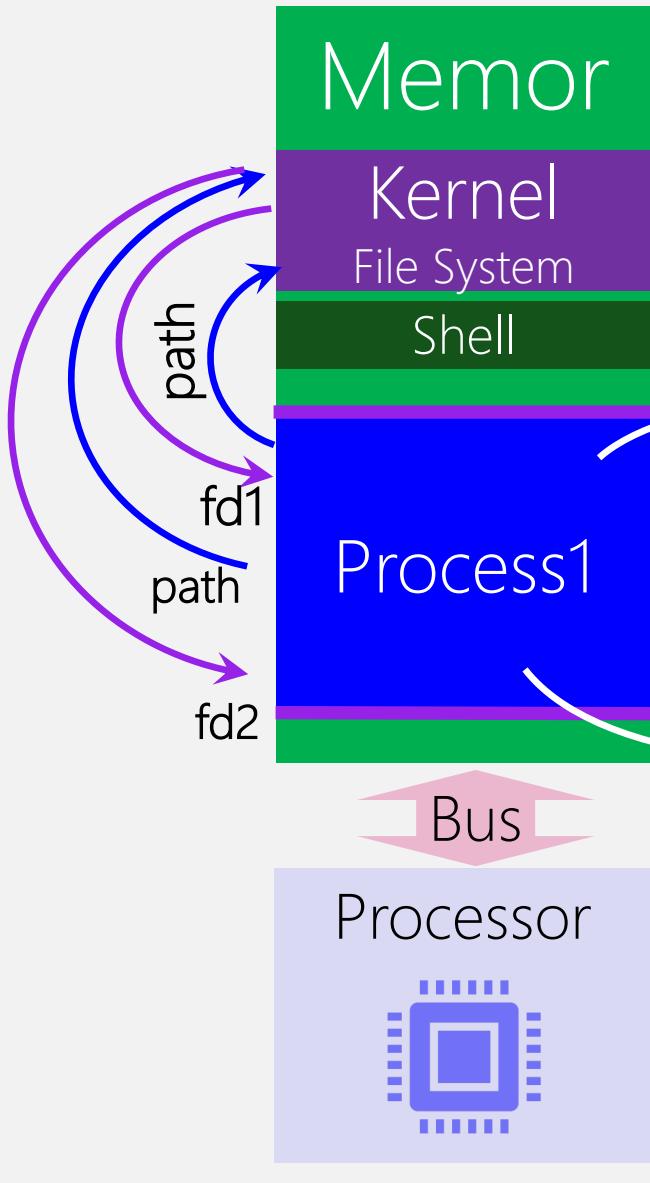
A Better User Case: Binary Search



dup and dup2

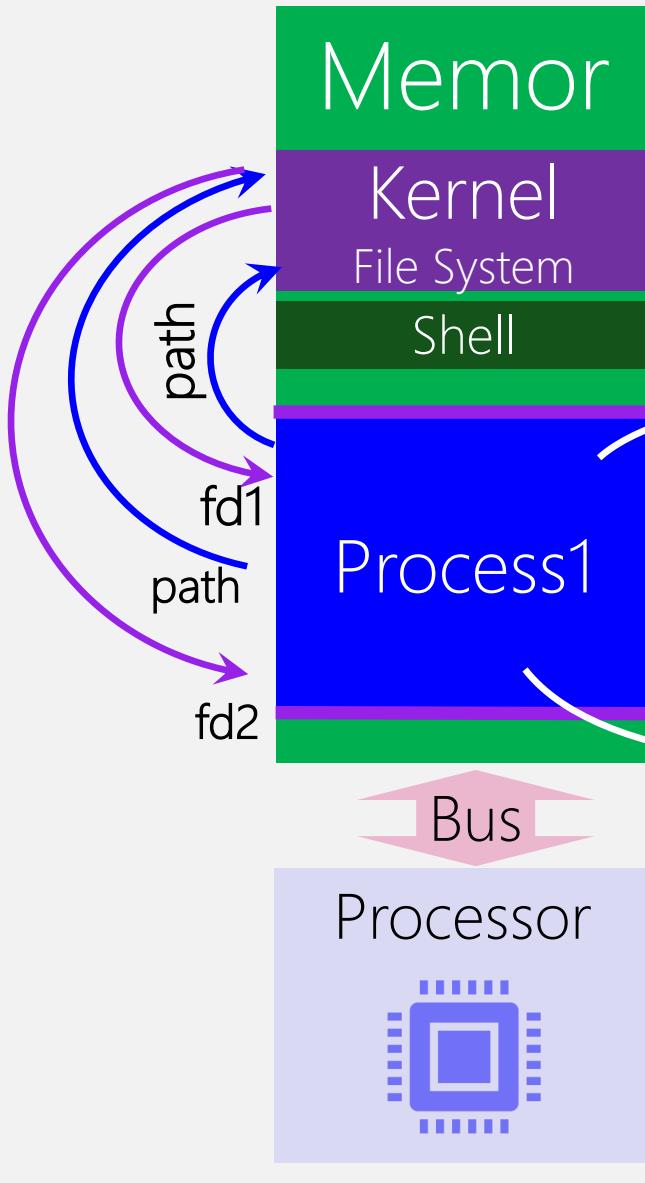
File System: High-Level

Computer



How to have multiple fds to the same file/device?

Computer



How to have multiple fds to the same file/device?
Multiple system calls for the same file!
`open()`

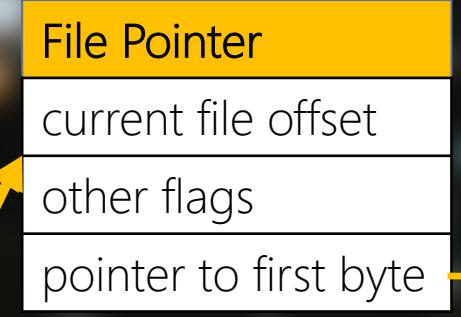
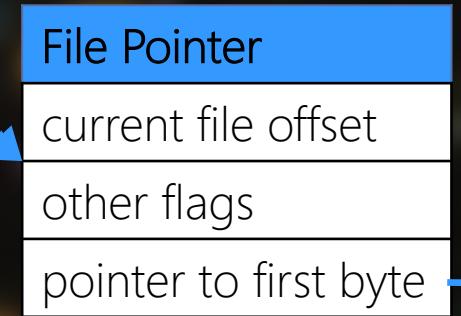
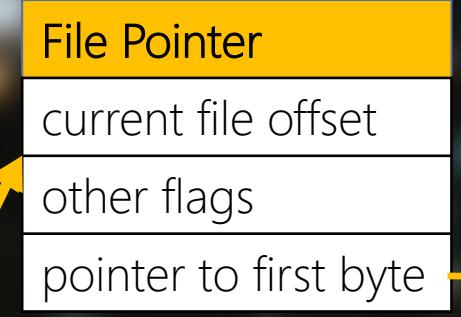
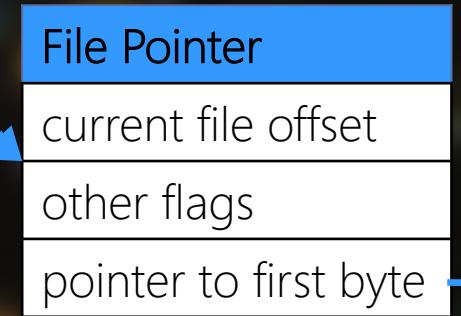
dup

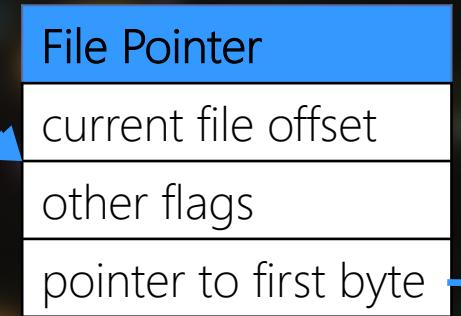
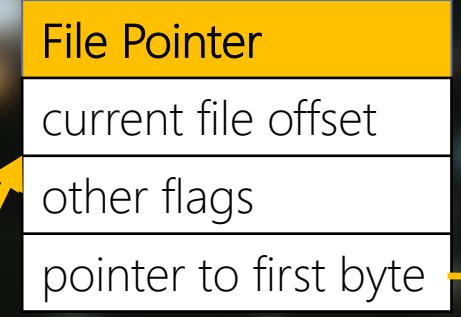
```
#include <unistd.h>
int dup(int fd);
```

new file descriptor to the same file if OK, -1 on error

Kernel File System

Process1

File Descriptors	File Pointer
fd1	
fd2	
fd3	
fd4	
...	



dup vs. multiple open

open: each returned fd has its own properties and current offset

```
fd1 = open("test.txt", O_RDONLY)
fd2 = open("test.txt", O_WRONLY)
```

dup vs. multiple open

dup: each returned fd has the same properties and current offset
fd2 = dup(fd1)

dup

Why do we duplicate fd?

I/O Redirection

STDIO \leftrightarrow File

STDERR \rightarrow STDOUT

Story

How STDIN_FILENO, STDOUT_FILENO, STDERR_FILENO are mapped and by who?

Computer

Memor

Kernel

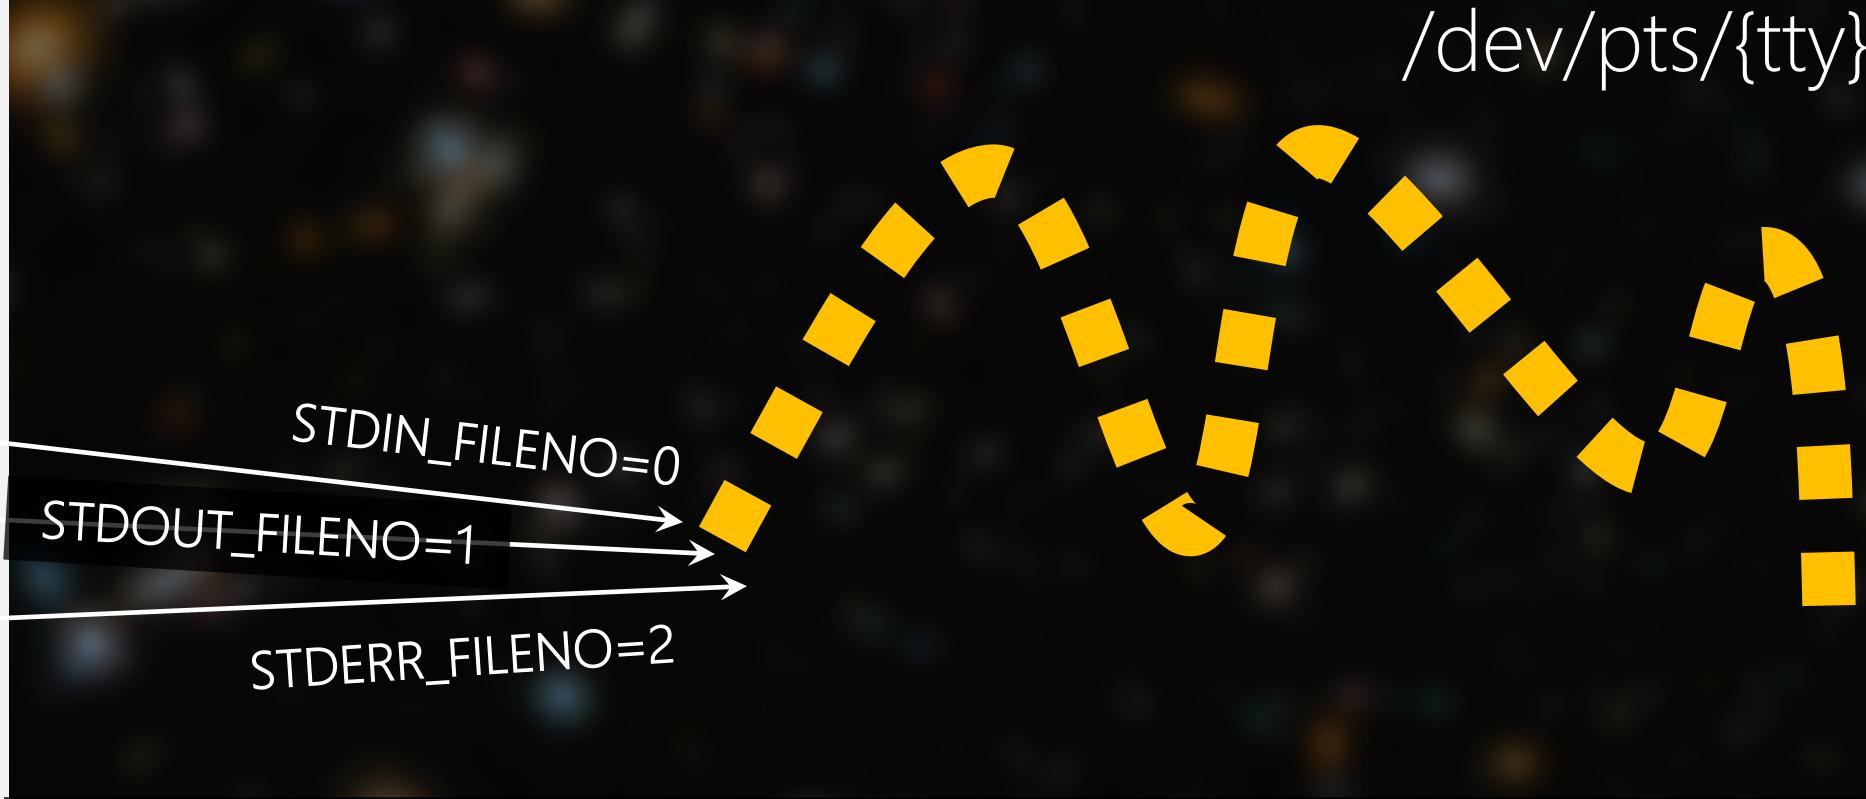
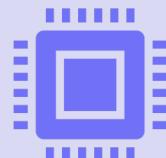
File System

Shell

Process1

Bus

Processor



When Shell bootstraps a program, it automatically opens three fds for the program (process):

`STDIN_FILENO = 0 : O_RDONLY`

`STDOUT_FILENO = 1 : O_WRONLY`

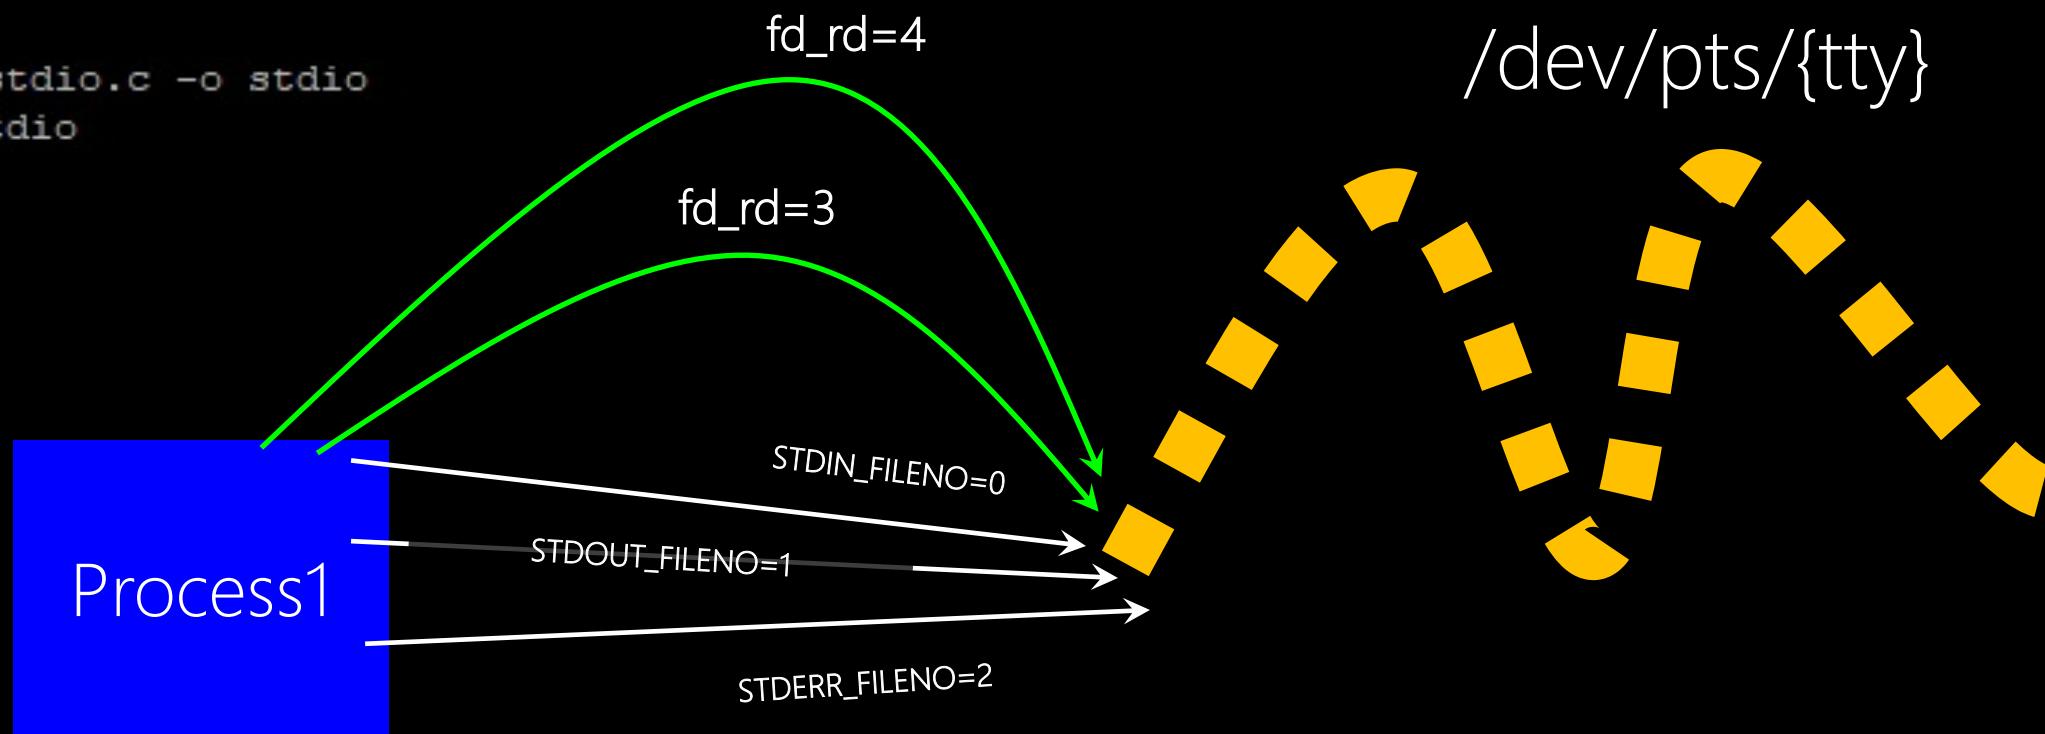
`STDERR_FILENO = 2 : O_WRONLY`

```
#include <fcntl.h>
#include <unistd.h>
void main(void) {
    char buf_rd[20];
    int fd_rd = open("/dev/pts/12", O_RDONLY);
    int res = read(fd_rd, buf_rd, 10);
    int fd_wr = open("/dev/pts/12", O_WRONLY);
    write(fd_wr, buf_rd, 10);
}
```

```
#include <fcntl.h>
#include <unistd.h>
void main(void) {
    char buf_rd[20];
    read(STDIN_FILENO, buf_rd, 10);
    write(STDOUT_FILENO, buf_rd, 10);
}
```

```
#include <fcntl.h>
#include <unistd.h>
void main(void) {
    char buf_rd[20];
    read(0, buf_rd, 10);
    write(1, buf_rd, 10);
}
```

```
hfani@charlie:~$ cc stdio.c -o stdio
hfani@charlie:~$ ./stdio
comp2560
comp2560
hfani@charlie:~$
```



Computer

Memor

Kernel

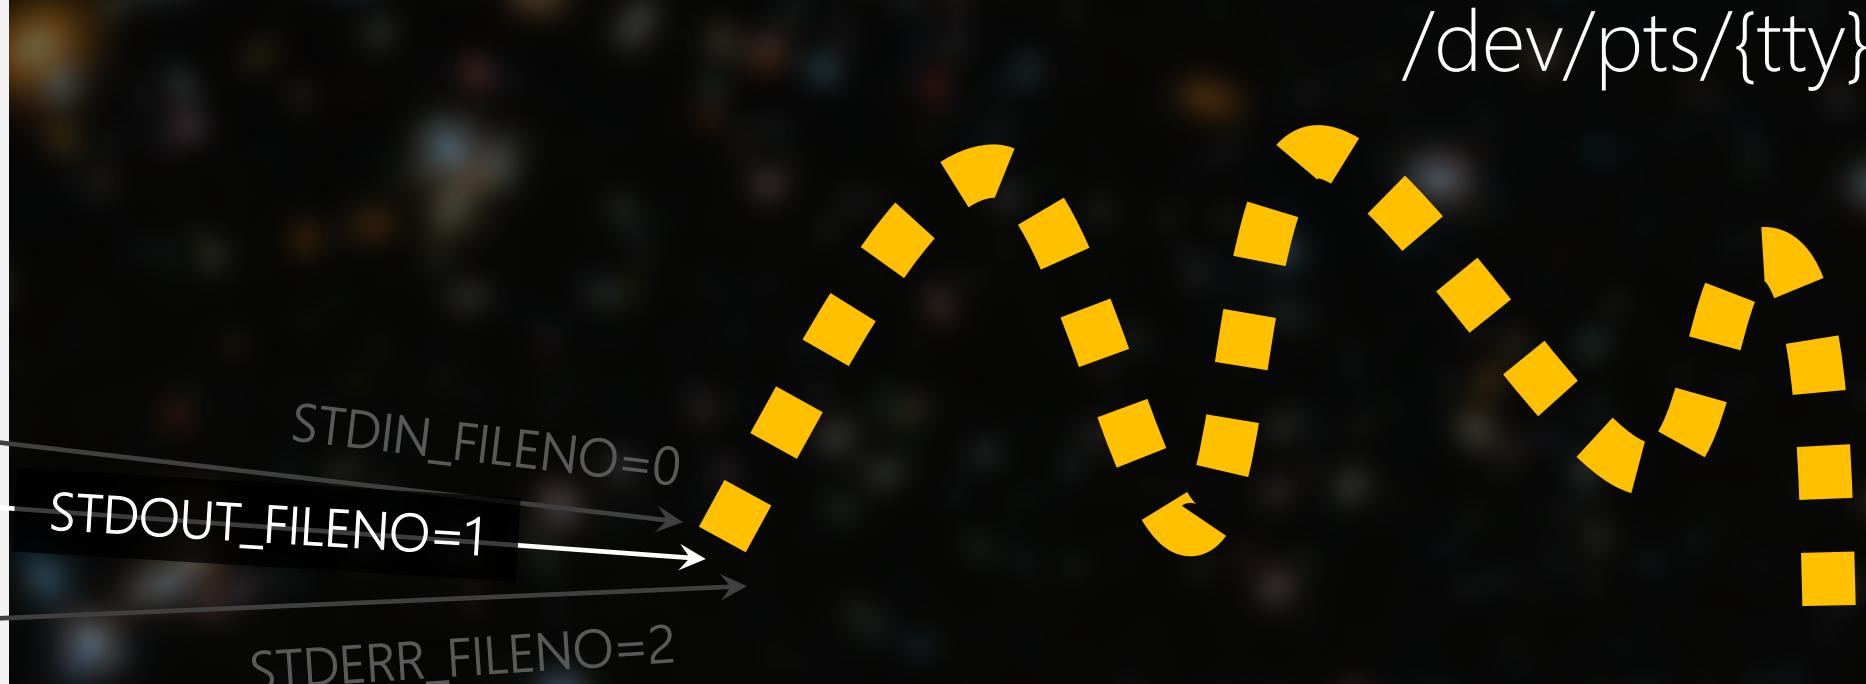
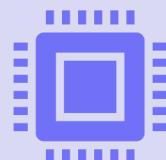
File System

Shell

Process1

Bus

Processor



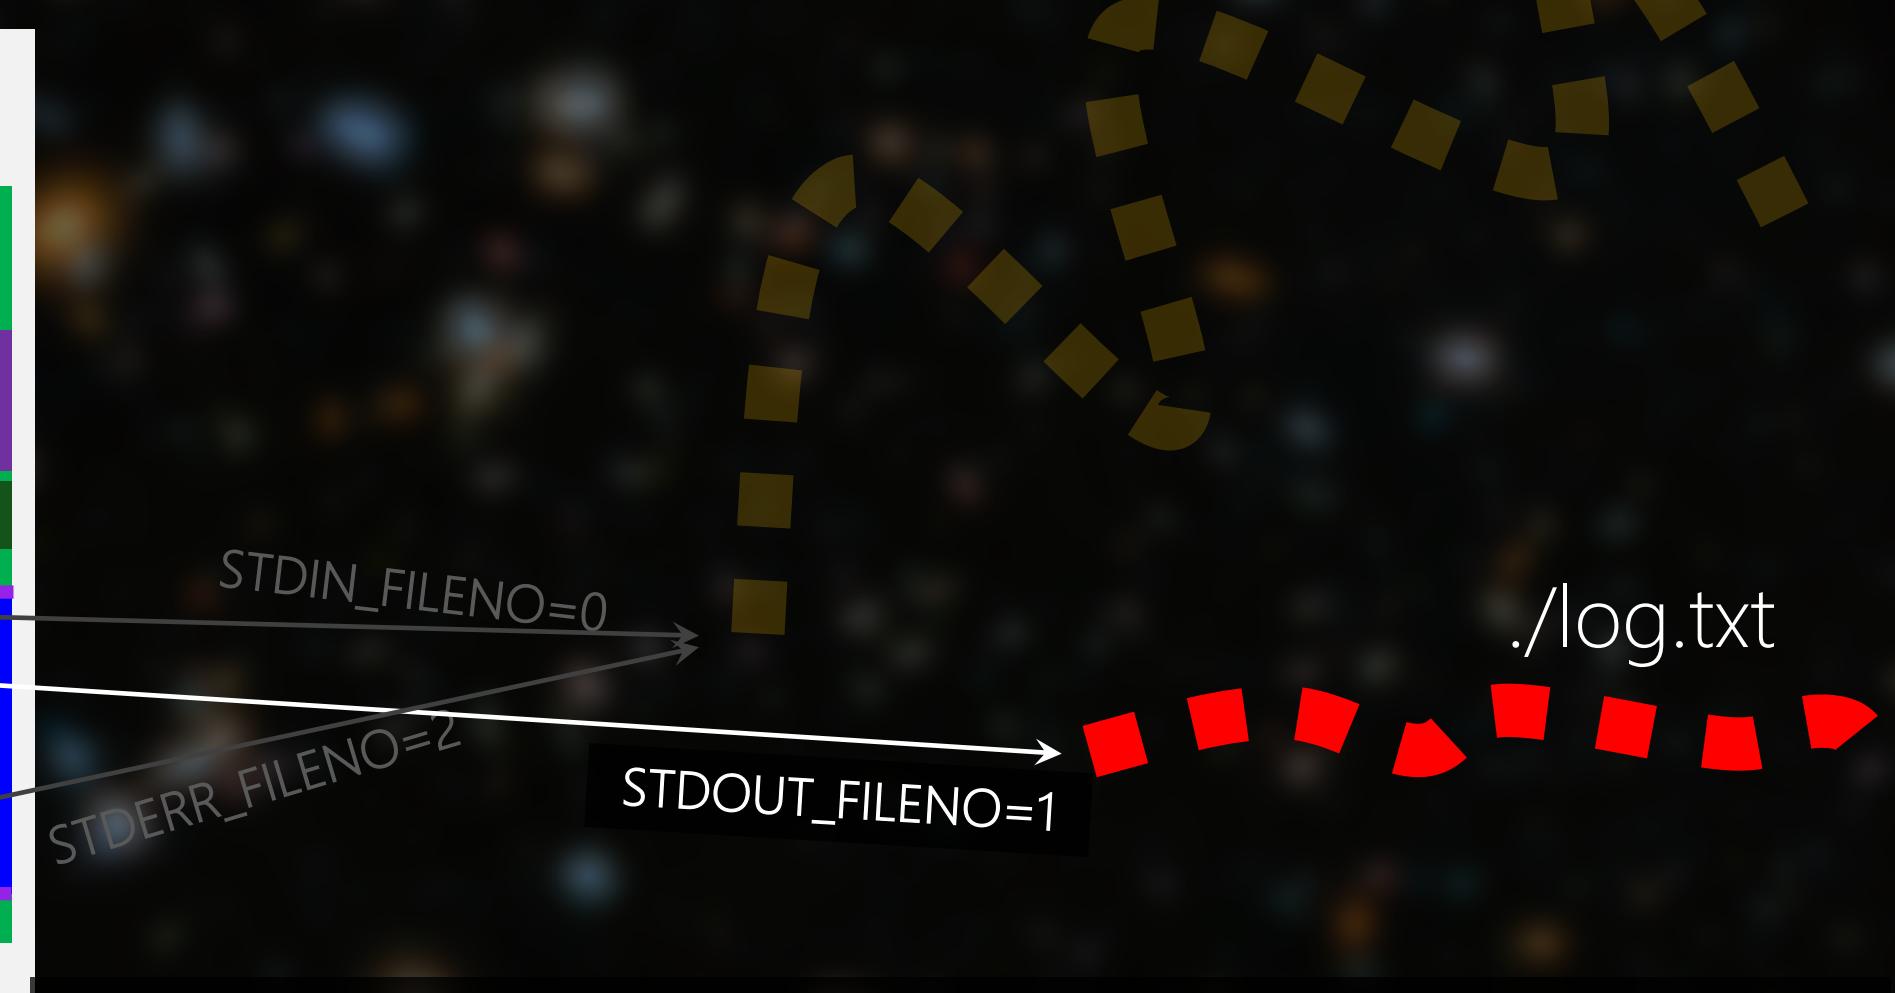
Now, we want to redirect output to a logging file `./log.txt`

`STDIN_FILENO = 0 : O_RDONLY`

`STDOUT_FILENO = 1 : O_WRONLY`

`STDERR_FILENO = 2 : O_WRONLY`

Computer



Now, we want to redirect output to a logging file `./log.txt`

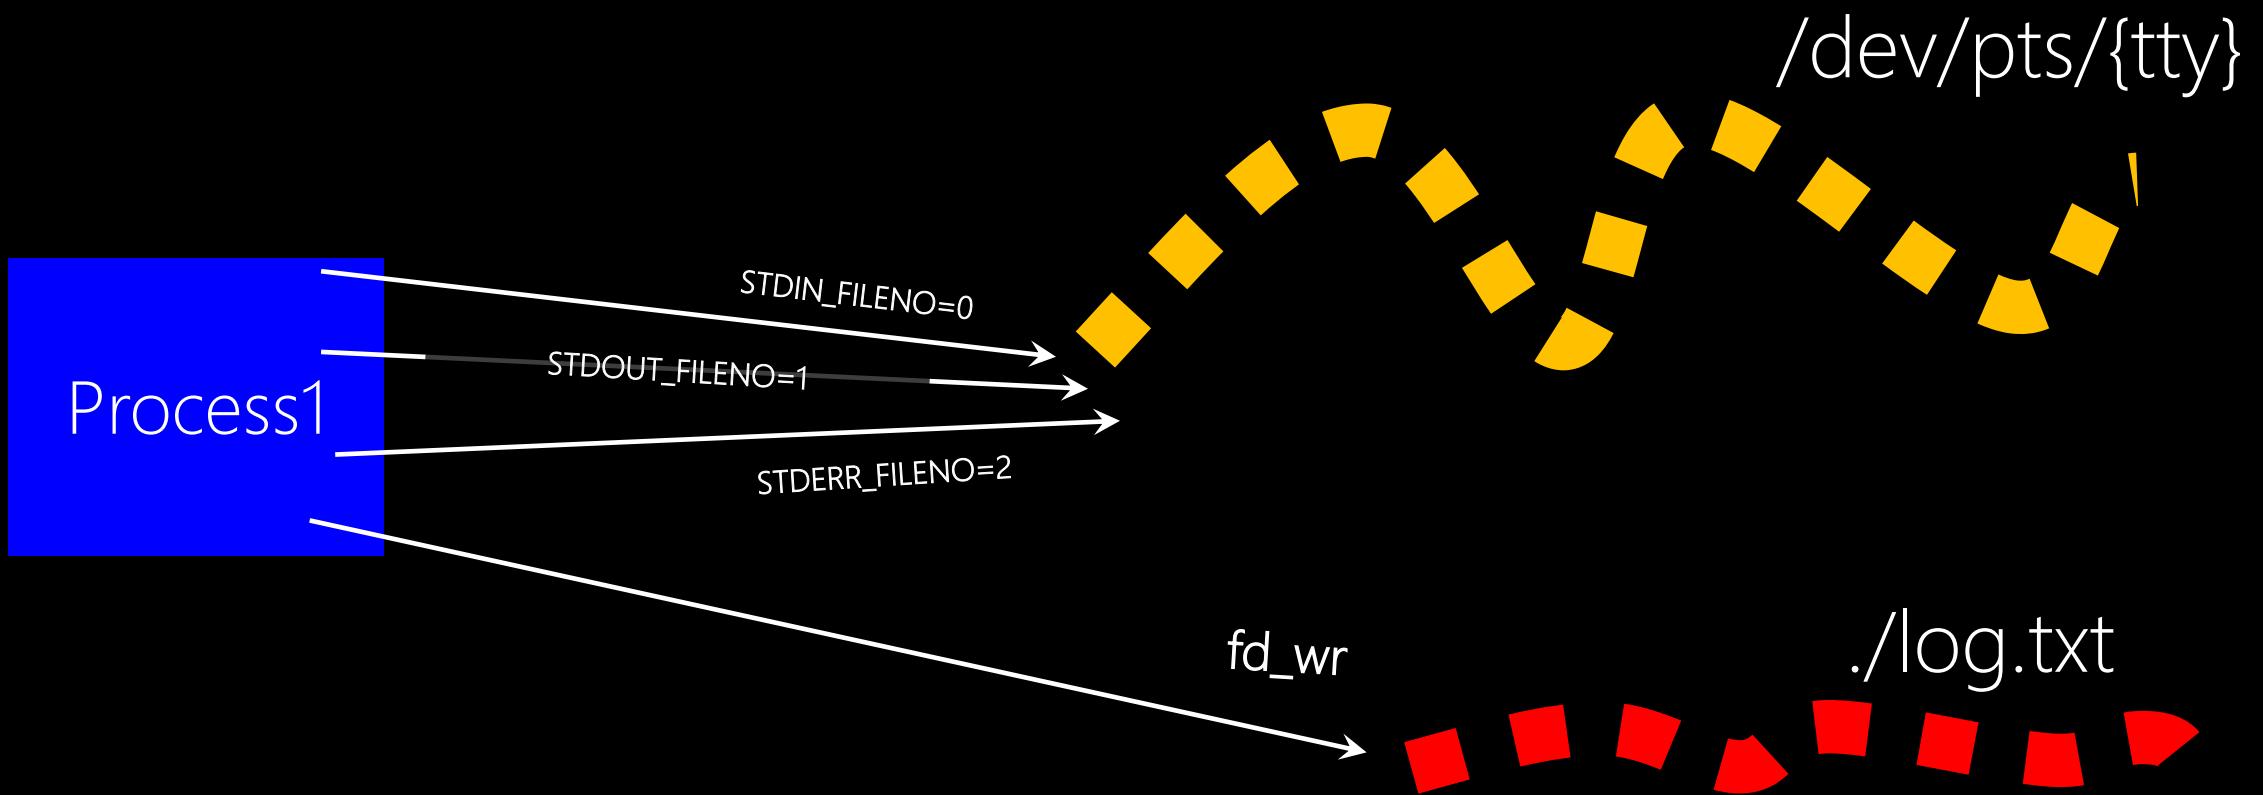
`STDIN_FILENO = 0 : O_RDONLY`

`STDOUT_FILENO = 1 : O_WRONLY`

`STDERR_FILENO = 2 : O_WRONLY`

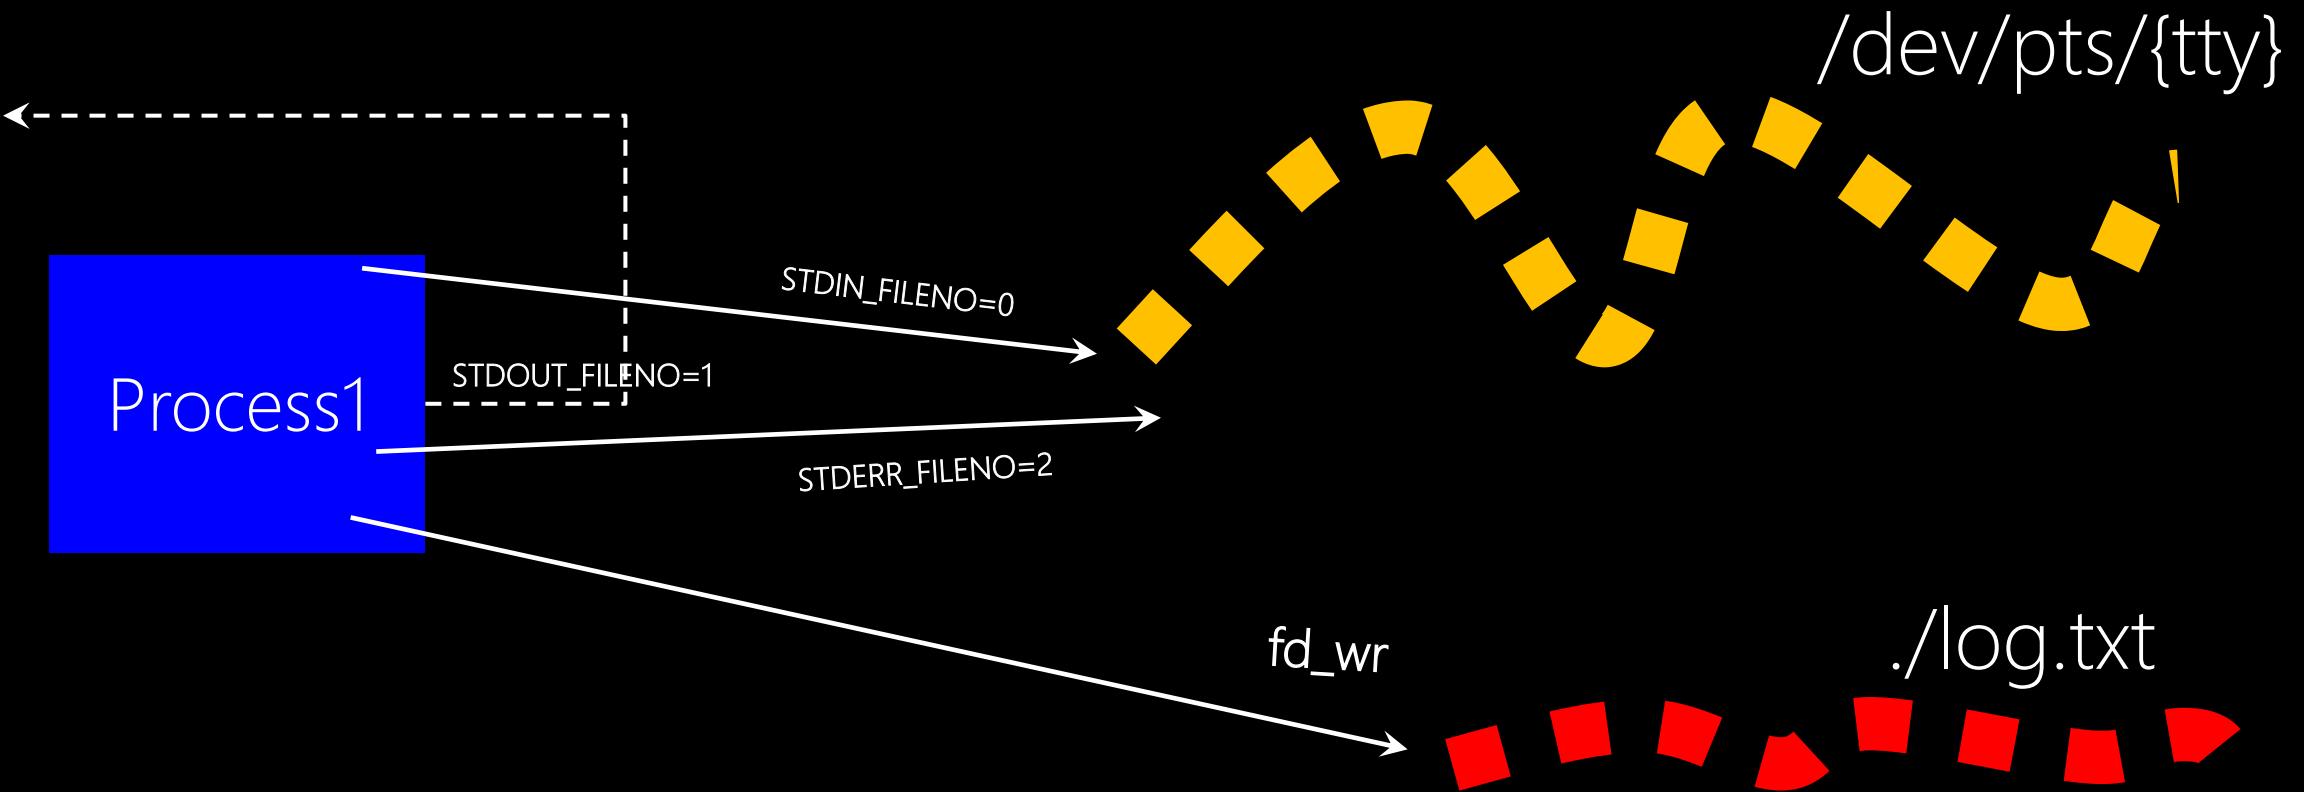
```
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
```



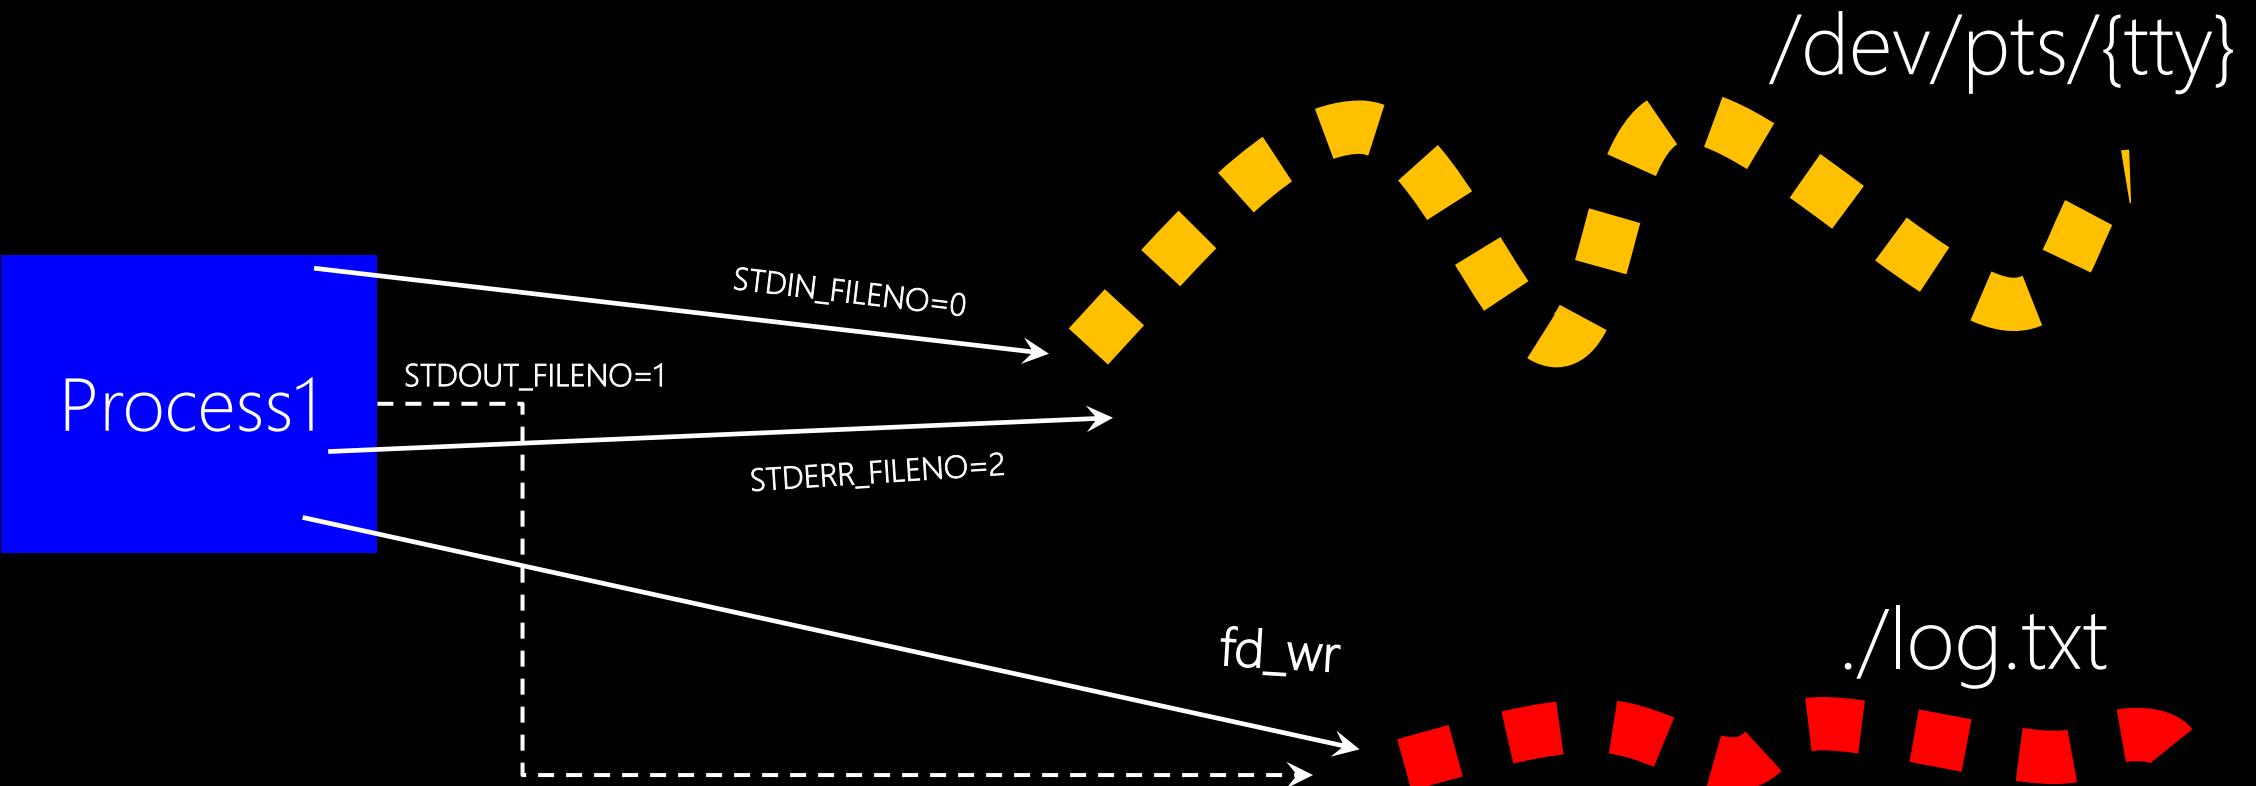
```
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
```



```
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 1
    //both fd_wr and new_fd are pointing to log.txt
```



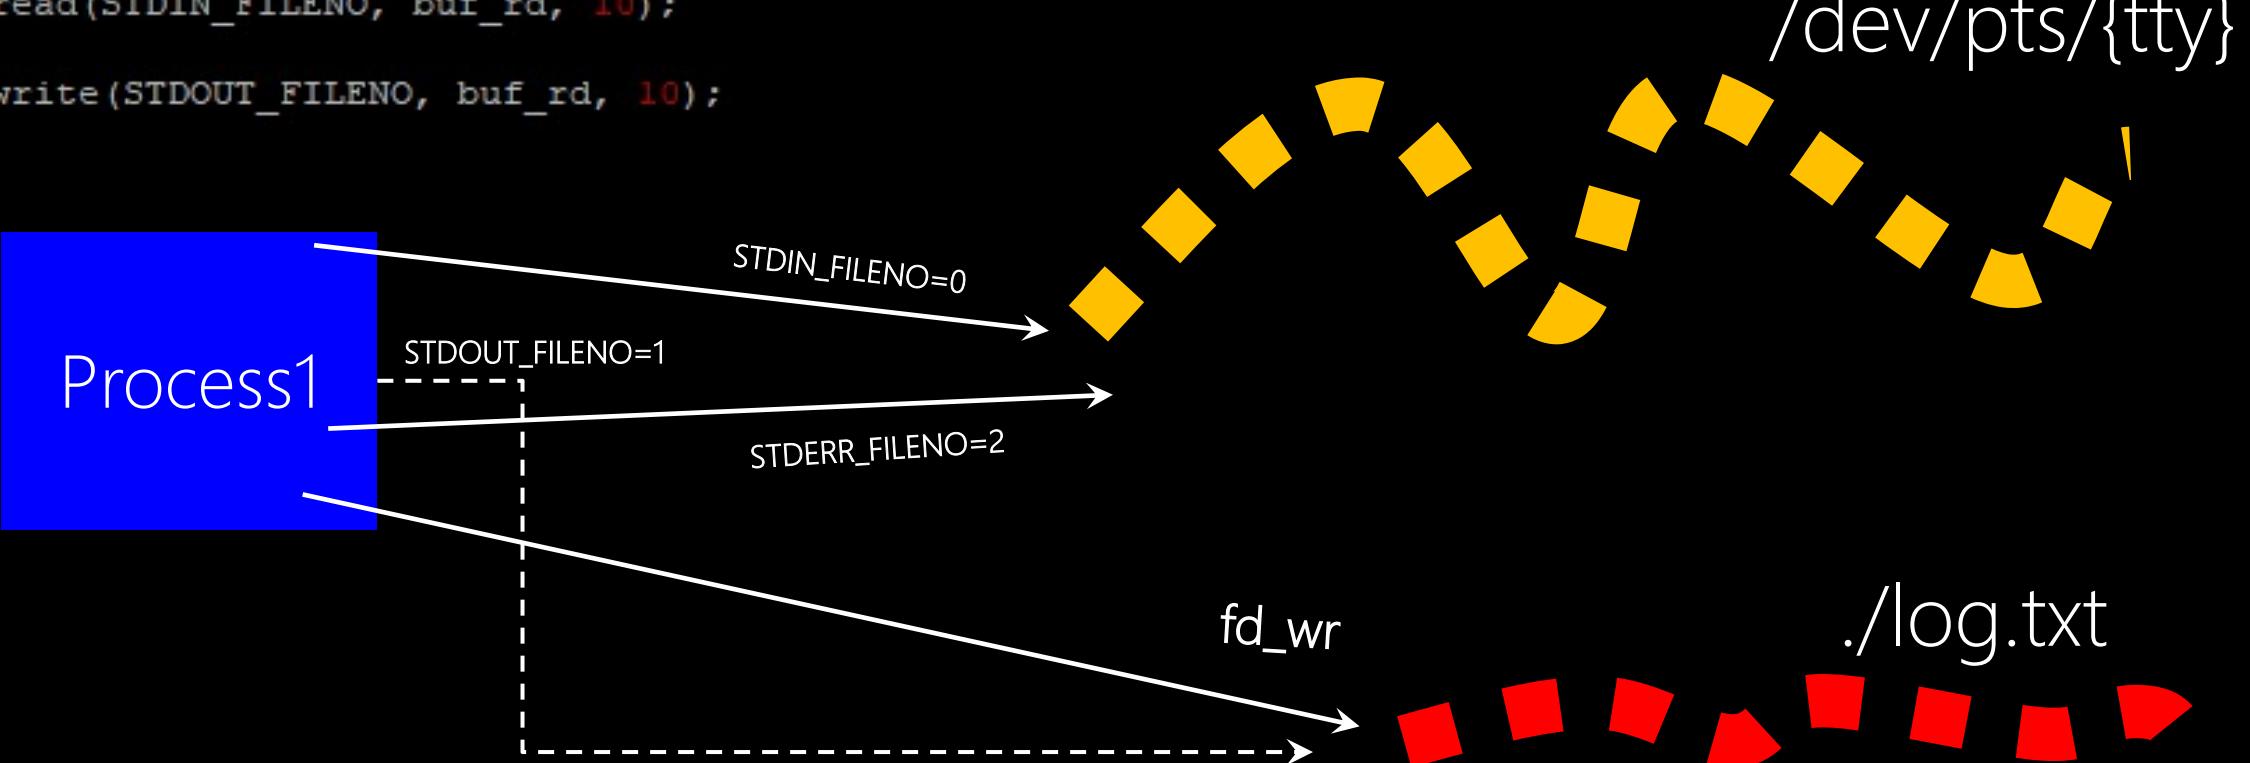
```

#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 1
    //both fd_wr and new_fd are pointing to log.txt

    read(STDIN_FILENO, buf_rd, 10);
    write(STDOUT_FILENO, buf_rd, 10);
}

```



```
hfani@charlie:~$ cc stdio_redirection.c -o stdio_redirection
```

```
hfani@charlie:~$ ./stdio_redirection
```

```
hey again!
```

```
hfani@charlie:~$
```

```
hfani@charlie:~$ vi log.txt
```

```
hey again!
```

```
~
```

Shell

\$./program > log.txt

We can ask the shell to do this redirection for us

{program file} > {new destination for STDOUT_FILENO}

Story

How about STDIN_FILENO? STDERR_FILENO?

Story

What does this mean and what is the benefit?

`fd = dup(0)`

dup2

At Home

I/O Efficiency buffered vs. unbuffered

`unistd.h` vs. `stdio.h`

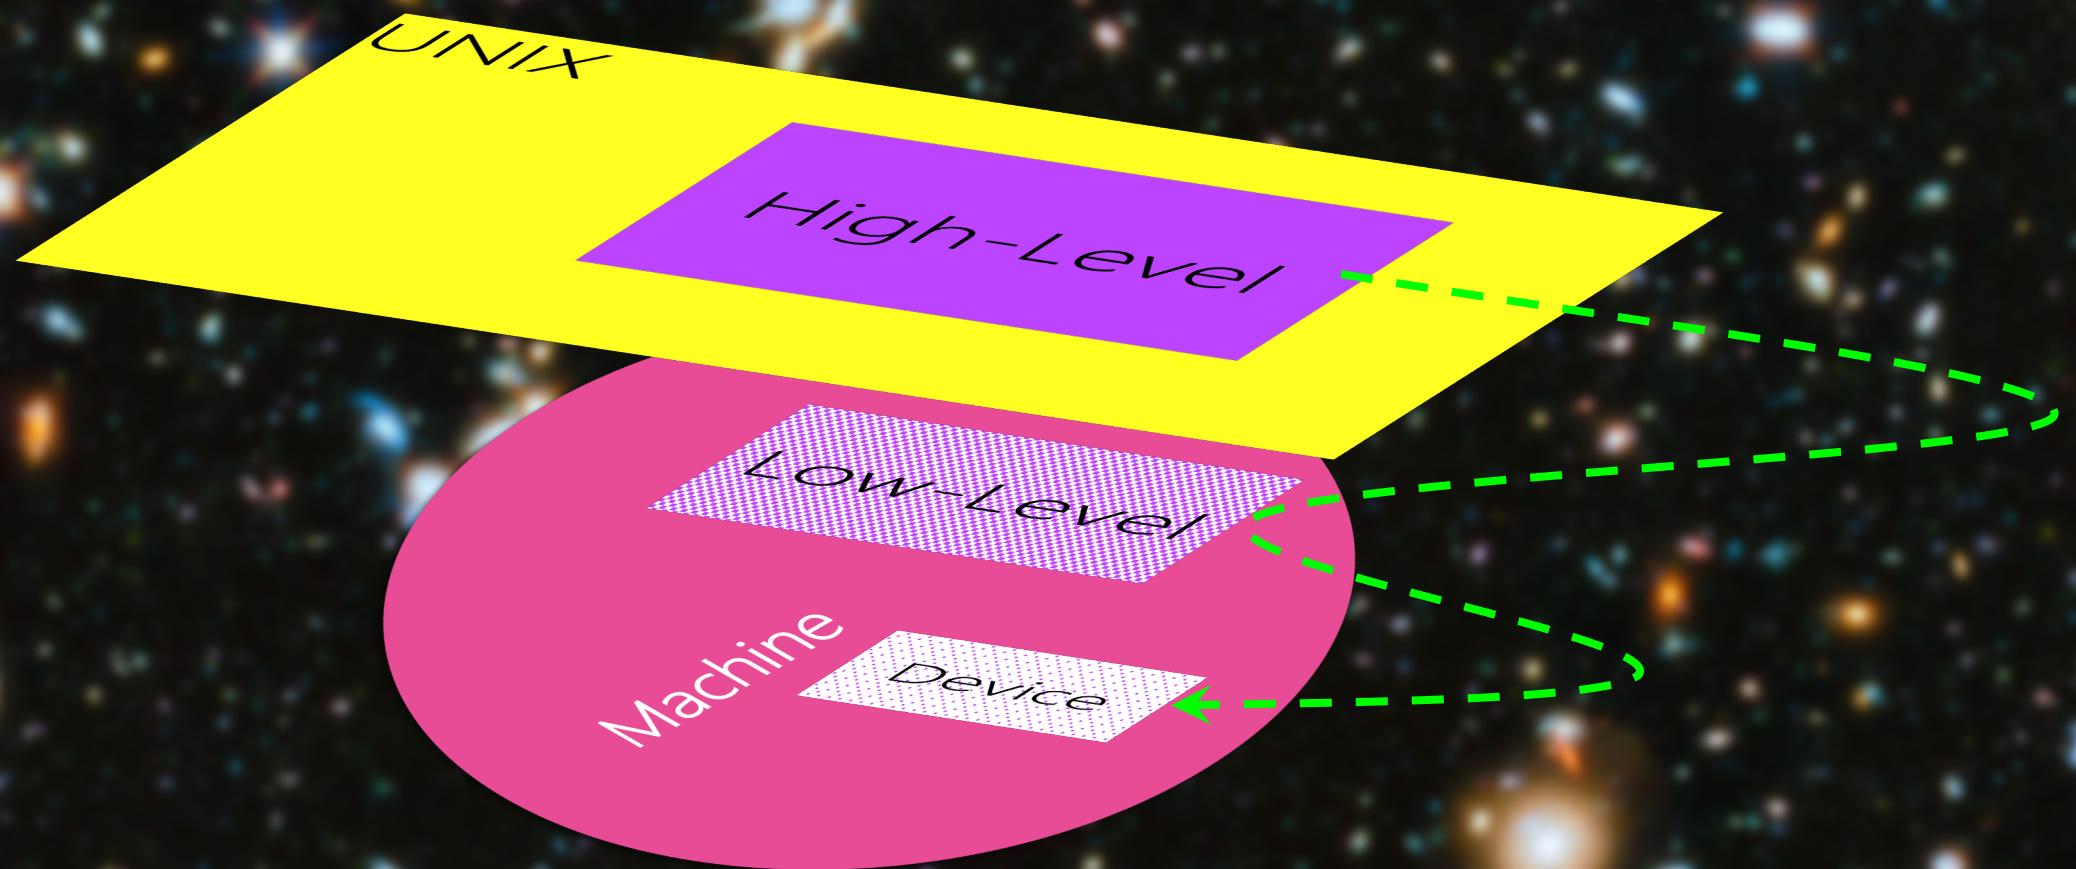
At Home.

File Sharing

Advanced! We won't cover it.

Atomic Operation

Advanced! We won't cover it.



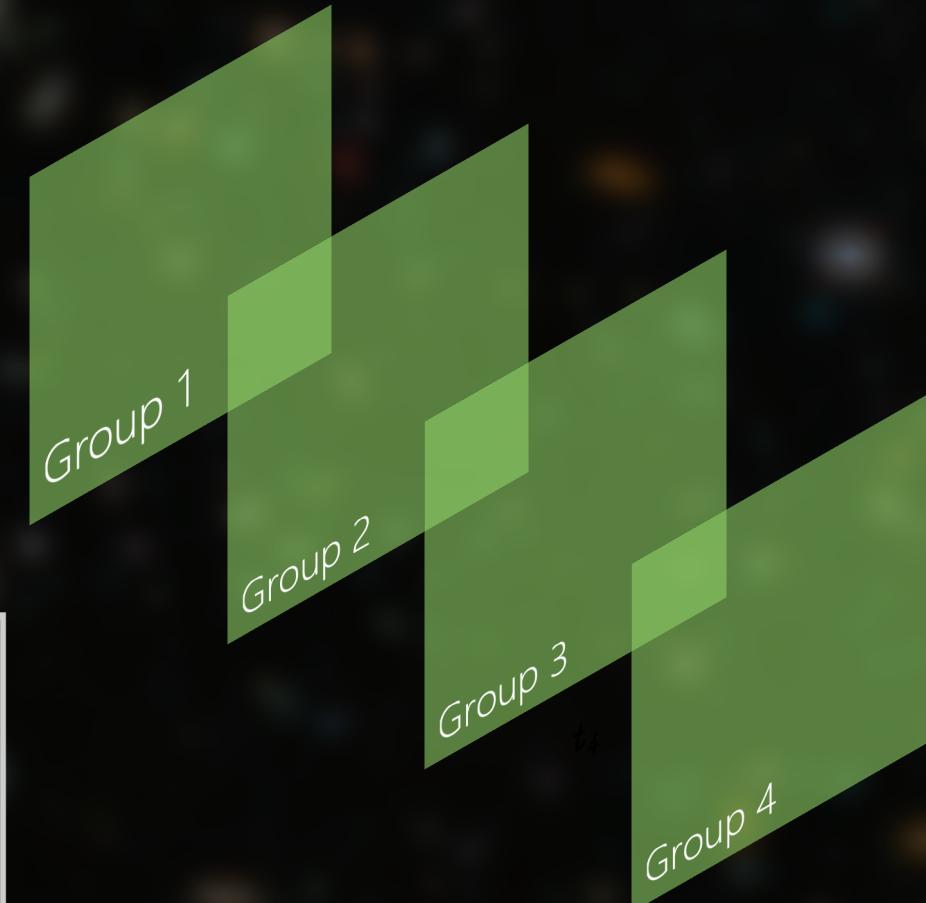
Storage File System

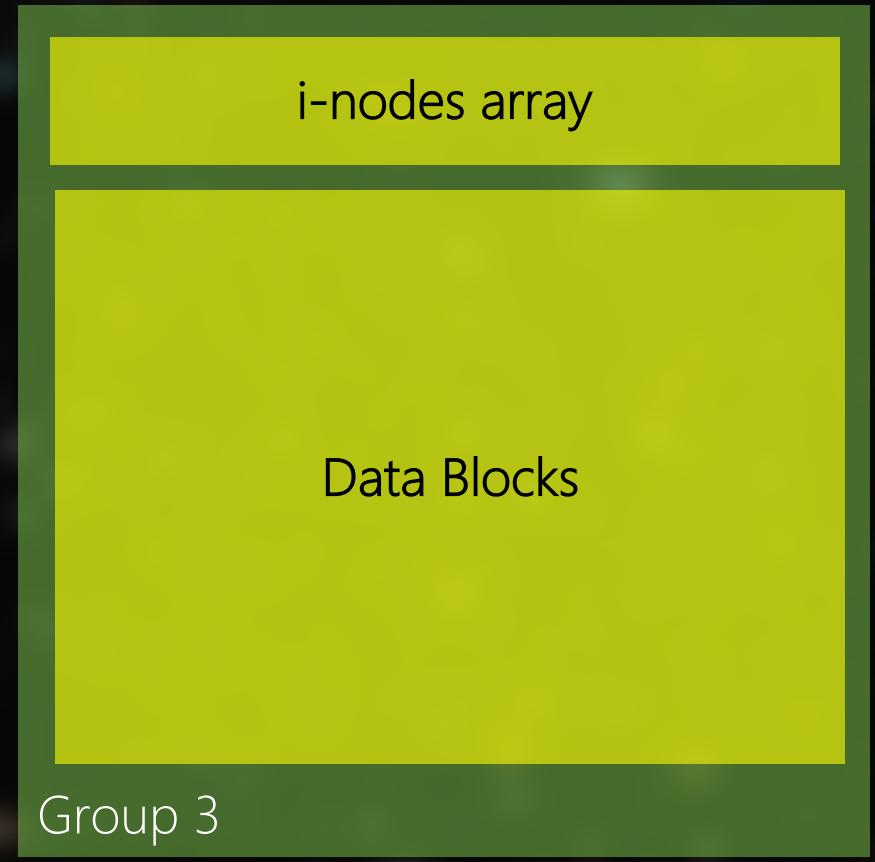
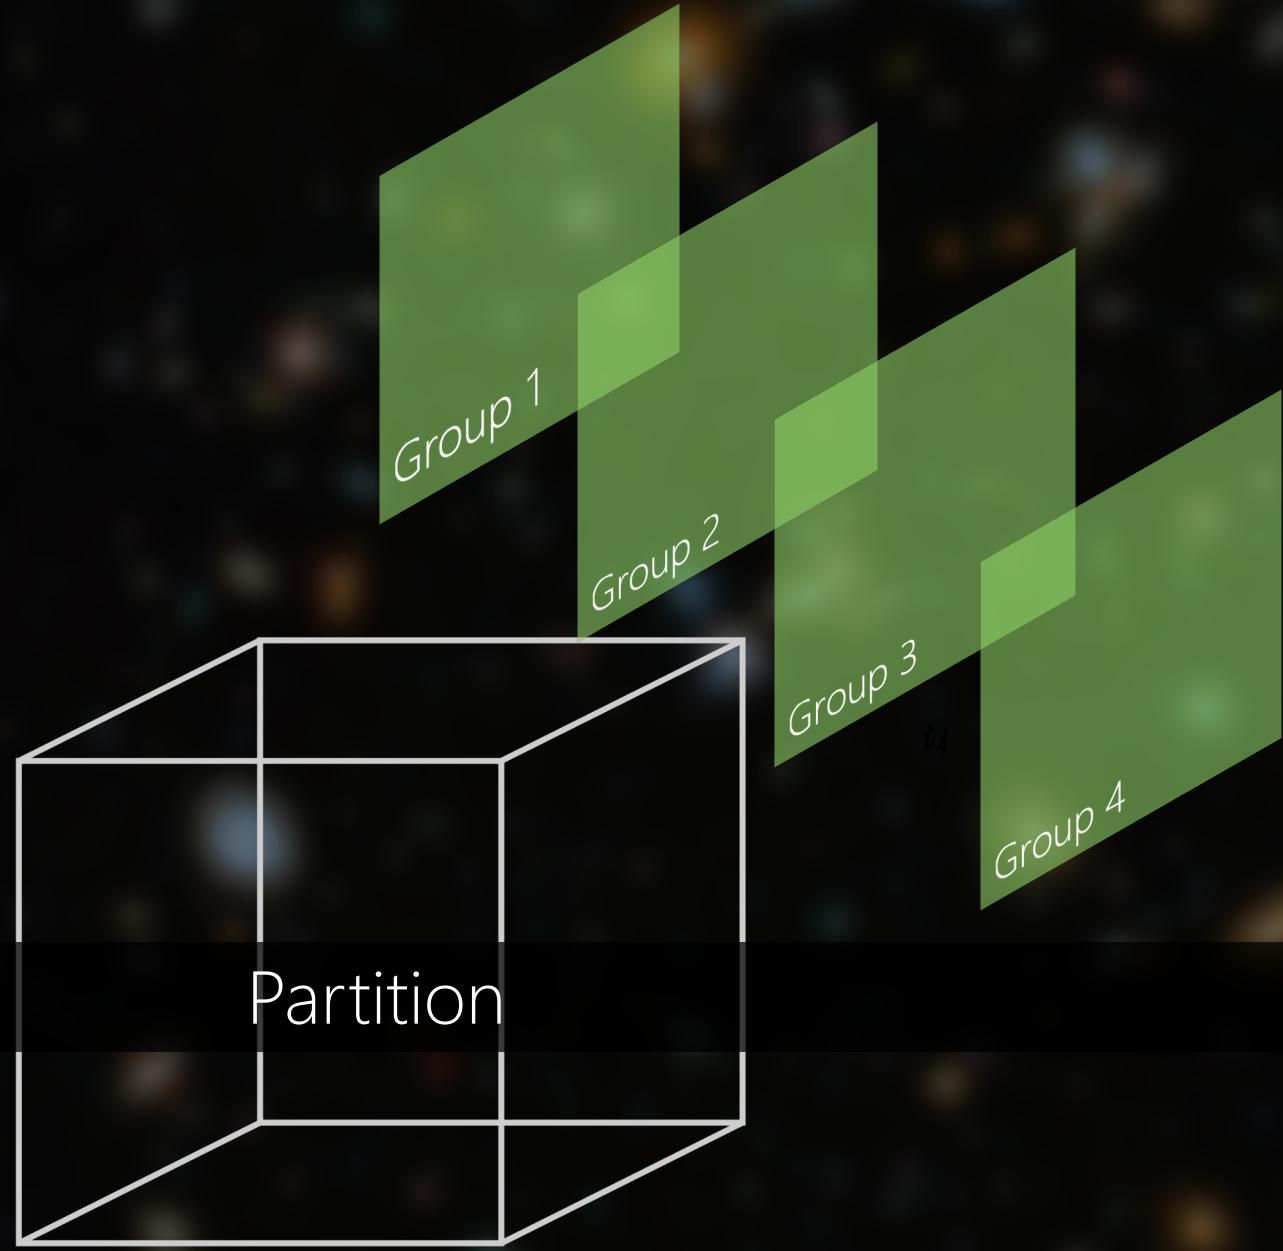
Disk File System

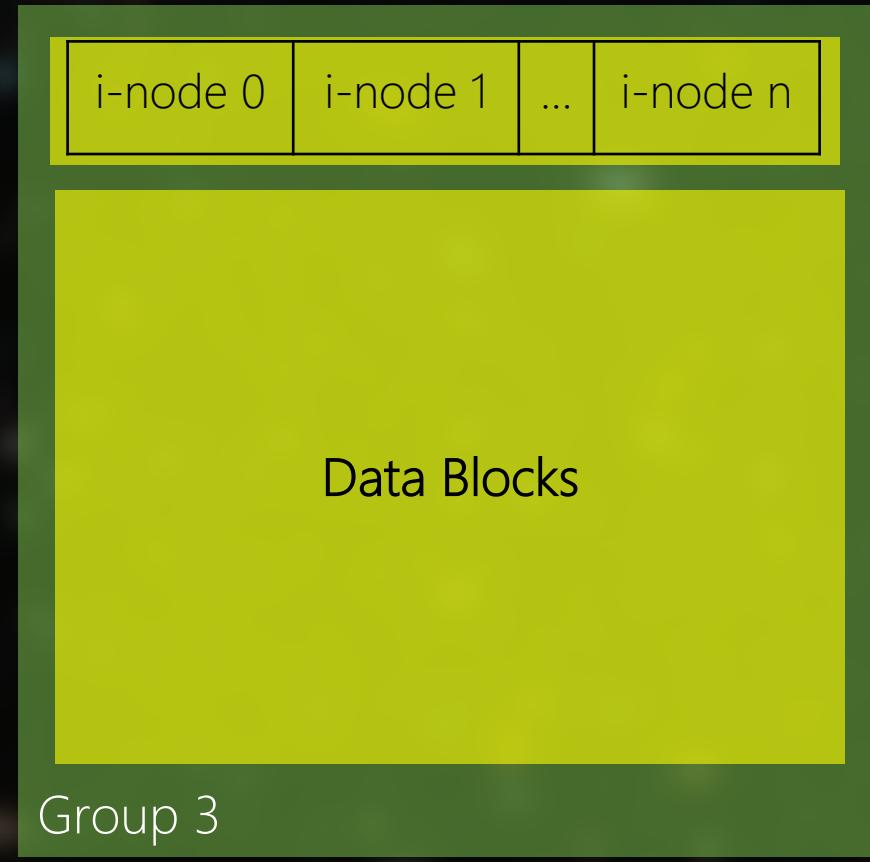
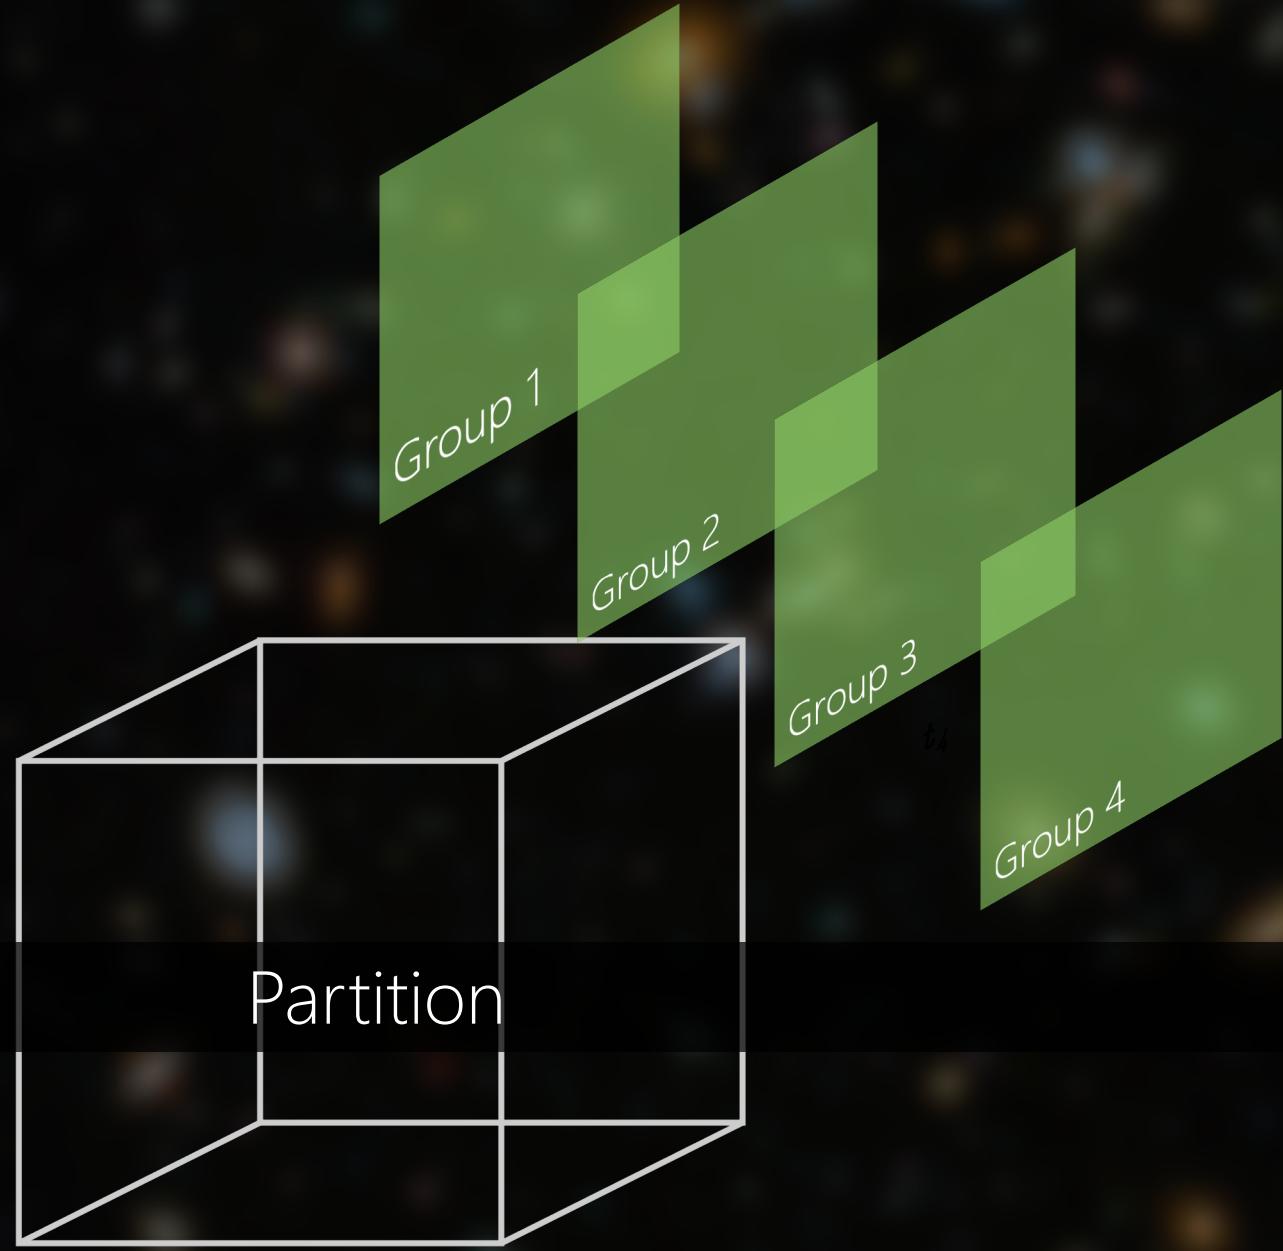
File System: Low Level







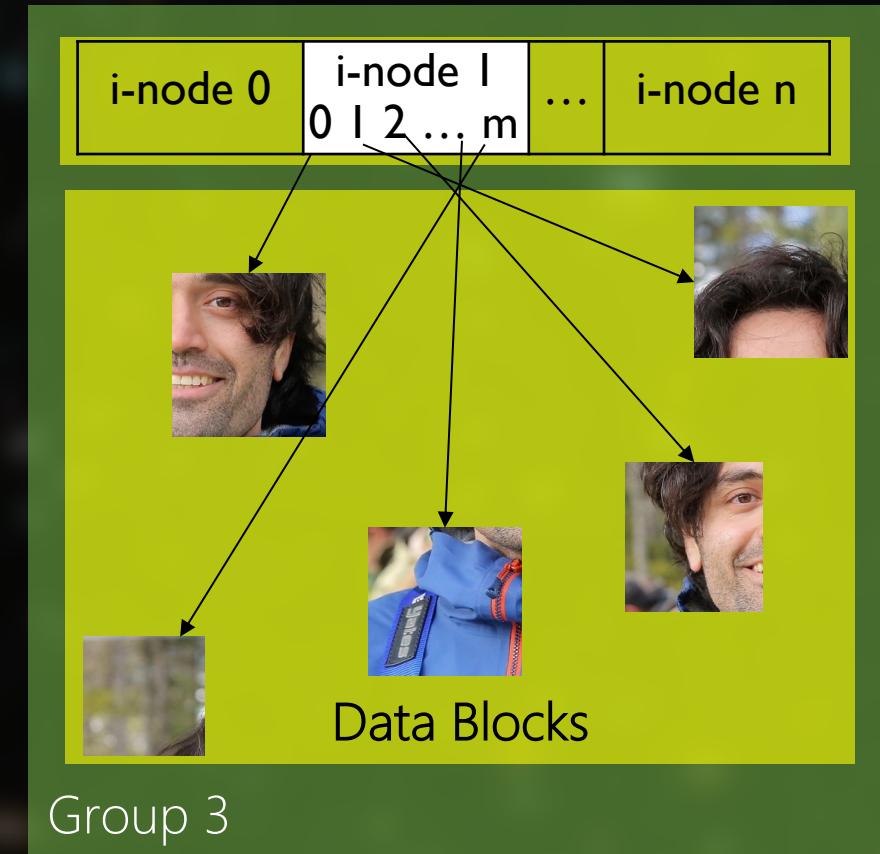
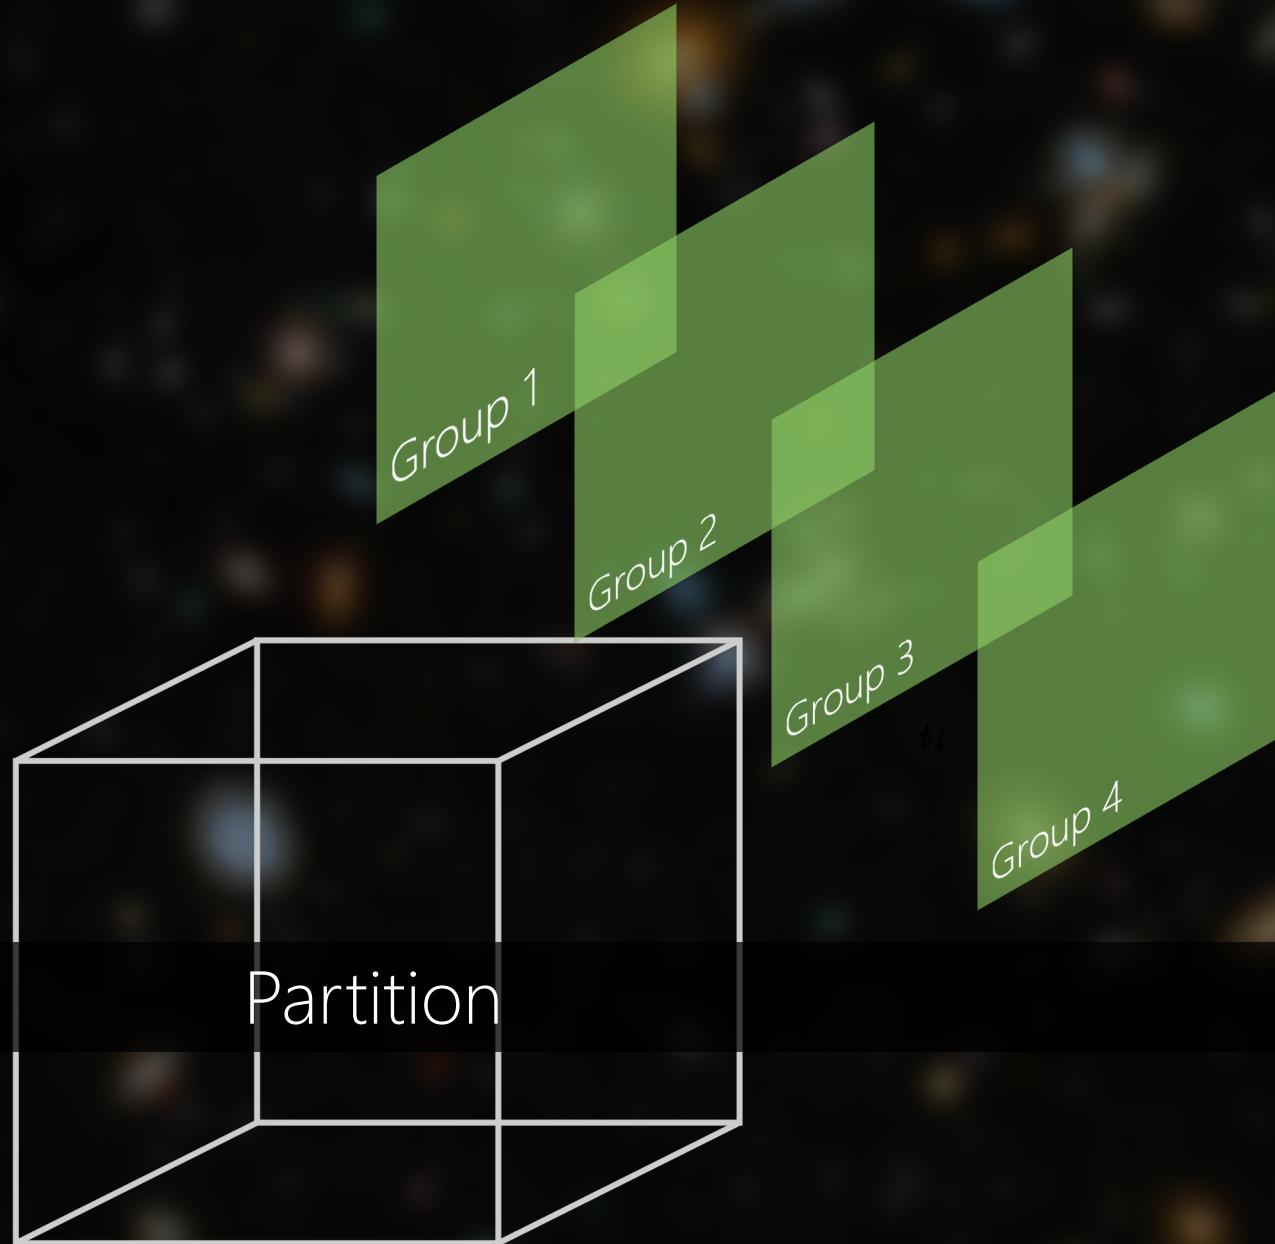




Each File has one i-node

```
$ ls -i {filename}
```

```
Administrator@hfani ~  
$ ls -i /cygdrive/c/hfani.jpeg  
1407374883666816 /cygdrive/c/hfani.jpeg
```



Each File has many data blocks

```
$ stat {filename}
```

```
Administrator@hfani ~
$ stat /cygdrive/c/hfani.jpeg
  File: /cygdrive/c/hfani.jpeg
  Size: 113327          Blocks: 112          IO Block: 65536 regular file
Device: 7ede359eh/2128491934d  Inode: 1407374883666816  Links: 1
Access: (0750/-rwxr-x---) Uid: ( 544/Administrators)  Gid: (197121/ None)
Access: 2021-10-25 00:07:34.744221300 -0400
Modify: 2021-07-05 22:48:41.000000000 -0400
```

/Path/Filename → i-node → Data Blocks

/cygdrive/c/hfani.jpeg → 1407374883666816 →



Directories

They are files. So, each of them has one i-node.
The content is not an image, audio, ... but mapping between filenames and their i-nodes

i-node 2 for root /

Data block b1

i-node#	filename
2	.
2	..
i1	home

i-node i1 for root '/home'

Data block b2

i-node#	filename
i1	.
2	..
i3	hfani

i-node i3 for root '/home/hfani'

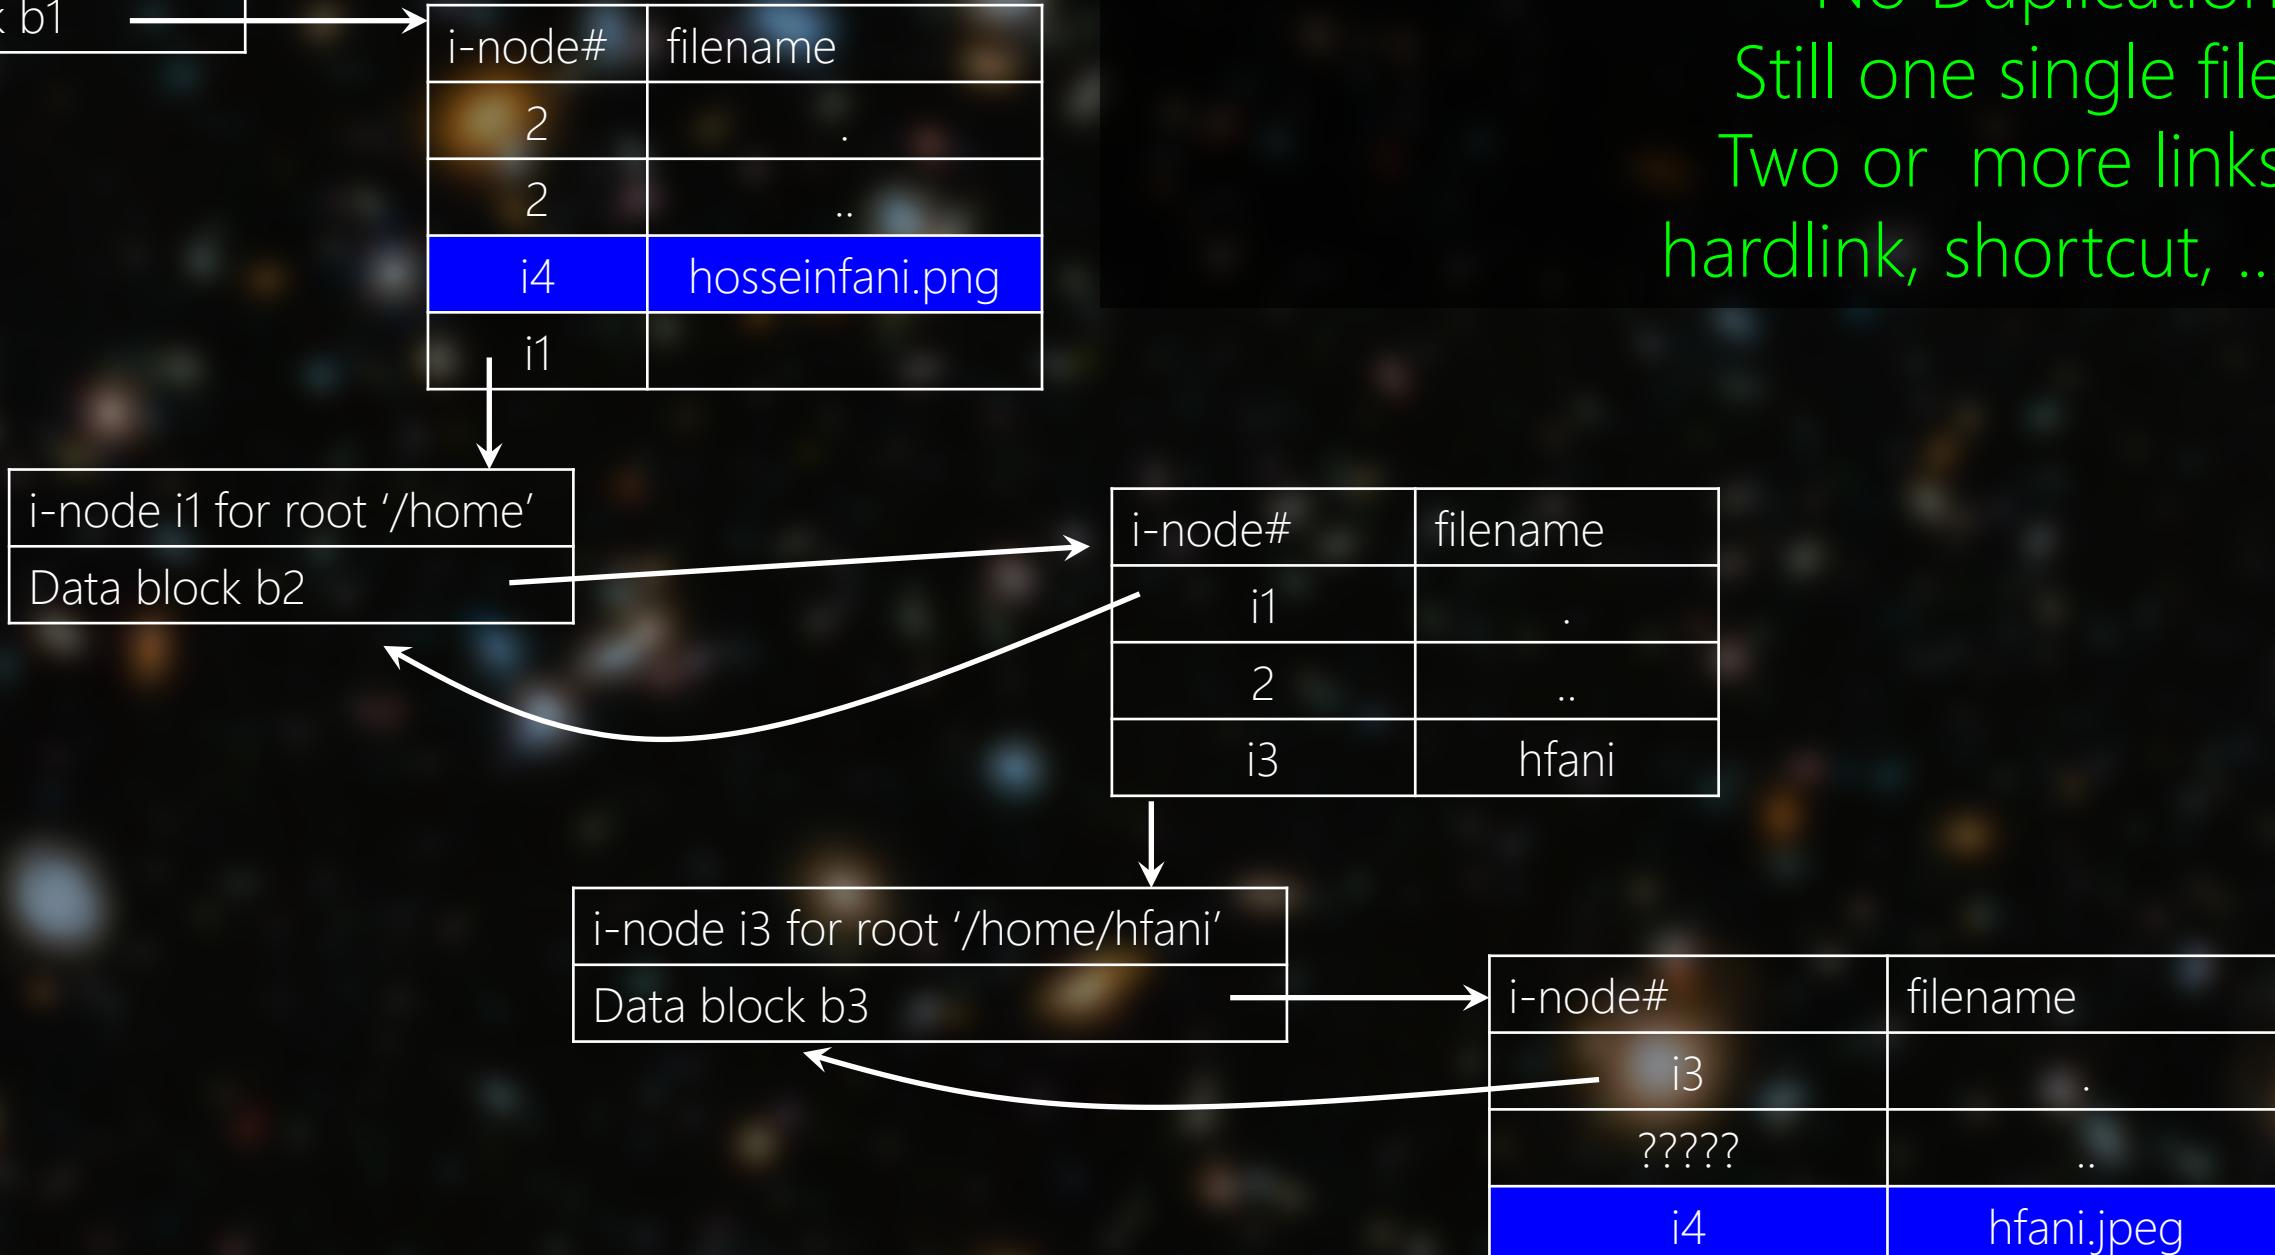
Data block b3

i-node#	filename
i3	.
?????	..
i4	hfani.jpeg

Directories

Is it possible to have multiple links to the same file? Yes.
How?

No Duplication!
Still one single file.
Two or more links.
hardlink, shortcut, ...



What happens when you **delete** a file?

What happens when you move a file?

What happens when you **copy** a file?
